

LAB - 07

Aim : Learn and apply the architecture of Apache Spark for data analytics. Solve Word Count program requirement using Apache Spark.

Description:

Spark is a robust and evolving engine for data analytics. Spark works with data using RDDs (Resilient Distributed Dataset) and data frames. The operations are immutable. The api allowed to process the data uses an approach of lazy processing. This allows spark to have various workflows checked out and optimize the processing. It uses graphs to represent the dependencies and operations, which are also utilized while optimizing.

Methodology:

There are two ways to run a program in apache spark

1. Using Spark REPL
2. By submitting a jar file with spark-submit

Spark shell:

Spark's shell provides a simple way to learn the API, as well as a powerful tool to analyze data interactively. It is available in either Scala (which runs on the Java VM and is thus a good way to use existing Java libraries) or Python.

● **Starting spark shell**

We can start a spark-shell by using spark-shell <deployment mode>

```
hadoop@hadoop-clone-21:~$ spark-shell --version
Welcome to

  ____      _
 / ___|  _ \| | | |
 \___ \| |_) | |_| |
  ___) | |_) | | | |
 |____|_|_|\___|_|_|_|

version 3.5.1

Using Scala version 2.12.18, Java HotSpot(TM) 64-Bit Server VM, 11
Branch HEAD
Compiled by user heartsavior on 2024-02-15T11:24:58Z
Revision fd86f85e181fc2dc0f50a096855acf83a6cc5d9c
Url https://github.com/apache/spark
Type --help for more information.
```

● **Spark deployment mode**

- I. Local mode

II. Standalone mode

III. Cluster mode(YARN)

❖ Working with Spark in Local mode

A. Running in local mode with maximum core as possible

```
spark-shell - -master "local[*]"
```

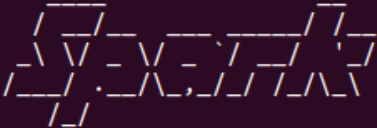
```
hadoop@hadoop-clone-12:~/Desktop$ spark-shell --master "local[*]"  
Setting default log level to "WARN".  
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).  
25/04/02 15:36:48 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable  
Spark context Web UI available at http://hadoop-clone-12:4040  
Spark context available as 'sc' (master = local[*], app id = local-1743588409616).  
Spark session available as 'spark'.  
Welcome to  
  
      / _ \   / _ \   / _ \   / _ \  
     / ___\ / ___\ / ___\ / ___\  
    /_____/_____/_____/_____\n\n        version 3.5.1
```

Using Scala version 2.12.18 (Java HotSpot(TM) 64-Bit Server VM, Java 11)
Type in expressions to have them evaluated.
Type :help for more information.

B. Running in local mode with 4 cores

```
spark-shell --master "local[4]"
```

```
hadoop@hadoop-clone-21:~$ spark-shell --master "local[4]"  
Setting default log level to "WARN".  
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).  
25/04/02 15:40:35 WARN NativeCodeLoader: Unable to load native-hadoop library fo  
r your platform... using builtin-java classes where applicable  
Spark context Web UI available at http://hadoop-clone-21:4040  
Spark context available as 'sc' (master = local[4], app id = local-1743588636434  
).  
Spark session available as 'spark'.  
Welcome to
```



version 3.5.1

```
Using Scala version 2.12.18 (Java HotSpot(TM) 64-Bit Server VM, Java 11)  
Type in expressions to have them evaluated.  
Type :help for more information.
```

```
scala>
```

❖ Wordcount program Using Spark shell REPL:

1. Create an RDD for inputfile

```
scala> val textFile = sc.textFile("file:///home/hadoop/inputdata.txt")
textFile: org.apache.spark.rdd.RDD[String] = file:///home/hadoop/inputdata.txt MapPartitionsRDD[9] at textFile at <console>:23
```

2. Display the data of the input file(optional)

```
scala> textFile.collect().foreach(println)
deer bear river
bear bear river
apple deer apple
bear river
deer apple apple cat
cat dog deer deer
penguin
```

Count number of lines

```
scala> val lineCount = textFile.count()
lineCount: Long = 7

scala> println(lineCount)
7
```

Perform word count

```
scala> val counts = textFile.flatMap(line => line.split(" "))
counts: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[10] at flatMap at <console>:23

scala>      .map(word => (word, 1))
res5: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[11] at map at <console>:24

scala>      .reduceByKey(_ + _)
res6: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[12] at reduceByKey at <console>:24
```

Display word count

```
scala> counts.collect().foreach(println)
deer
bear
river
bear
bear
river
apple
dear
apple
bear
river
dear
apple
apple
cat
cat
dog
dear
deer
penguin
```

Collect the output

```
scala> inputfile.collect()
res12: Array[String] = Array(deer bear river, bear bear river, apple
dear apple, bear river, dear apple apple cat, cat dog dear deer, peng
uin)
```

3. Create a mapper and reducer

```
scala> val counts=inputfile.flatMap(line=>line.split("")).map(word=>(
word,1)).reduceByKey(_+_);
counts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[18] at
reduceByKey at <console>:23
```

● RDD.toDebugString()

```
scala> counts.toDebugString
res13: String =
(2) ShuffledRDD[18] at reduceByKey at <console>:23 []
+- (2) MapPartitionsRDD[17] at map at <console>:23 []
|   MapPartitionsRDD[16] at flatMap at <console>:23 []
|   file:///home/hadoop/inputdata.txt MapPartitionsRDD[15] at text
File at <console>:23 []
|   file:///home/hadoop/inputdata.txt HadoopRDD[14] at textFile at
<console>:23 []
```

- **What is RDD Persistence and caching**

Spark RDD persistence is an optimization technique in which saves the result of RDD evaluation. Using this we save the intermediate result so that we can use it further if required. It reduces the computation overhead.

cache() and persist()

RDD.cache()

Persist this RDD with the default storage level (MEMORY_ONLY).

```
scala> counts.cache()  
res14: counts.type = ShuffledRDD[18] at reduceByKey at <console>:23
```

4. Action - perform actions on created RDD

- collect -


Counts.collect will collect the data and display it on console

```
scala> counts.collect  
res15: Array[(String, Int)] = Array((d,6), (p,9), (t,2), (b,4), (" ",  
13), (n,2), (v,3), (r,15), (l,4), (e,19), (a,13), (i,4), (u,1), (o,1)  
, (g,2), (c,2))
```

- If you want to save your data as output file we can use the following command

```
scala> counts.saveAsTextFile("file:///home/hadoop/output_WC");
```

This is created in /home/hadoop



part-00000

- 1 (d,6)
- 2 (p,9)
- 3 (t,2)
- 4 (b,4)
- 5 (,13)
- 6 (n,2)
- 7 (v,3)
- 8 (r,15)
- 9 (l,4)

part-00001

- 1 (e,19)
- 2 (a,13)
- 3 (i,4)
- 4 (u,1)
- 5 (o,1)
- 6 (g,2)
- 7 (c,2)

Observing DAG visualization

← → ↺ localhost:4040/jobs/ ☆ 📄 ⬇️ 🔴 Finish update

Spark 3.5.1 Jobs Stages Storage Environment Executors Spark shell application UI

Spark Jobs (?)

User: hadoop
Total Uptime: 24 min
Scheduling Mode: FIFO
Completed Jobs: 6

▶ Event Timeline

▼ Completed Jobs (6)


Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

Job Id ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
5	runJob at SparkHadoopWriter.scala:83 runJob at SparkHadoopWriter.scala:83	2025/04/02 16:04:52	0.1 s	1/1 (1 skipped)	2/2 (2 skipped)
4	collect at <console>:24 collect at <console>:24	2025/04/02 16:03:43	0.1 s	2/2	4/4
3	collect at <console>:24 collect at <console>:24	2025/04/02 15:57:13	16 ms	1/1	2/2
2	collect at <console>:24 collect at <console>:24	2025/04/02 15:52:50	19 ms	1/1	2/2
1	count at <console>:23 count at <console>:23	2025/04/02 15:51:55	24 ms	1/1	2/2
0	collect at <console>:24 collect at <console>:24	2025/04/02 15:50:29	0.2 s	1/1	2/2

Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

Go to timeline of events and explore DAG visualization for the submitted job:

Details for Job 5

Status: SUCCEEDED**Submitted:** 2025/04/02 16:04:52**Duration:** 0.1 s**Completed Stages:** 1**Skipped Stages:** 1[▶ Event Timeline](#)[▼ DAG Visualization](#)

We can also observe the status of the jobs as below:

[▼ Completed Stages \(1\)](#)

Page: 1

Stage Id ▼	Description	Submitted	Duration	Tasks: Succeeded/Total
7	runJob at SparkHadoopWriter.scala:83 +details	2025/04/02 16:04:52	90 ms	2/2

Page: 1

[▼ Skipped Stages \(1\)](#)

Page: 1

Stage Id ▼	Description	Submitted	Duration	Tasks: Succeeded/Total	Input
6	map at <console>:23 +details	Unknown	Unknown	0/2	

Page: 1

Wordcount Program by taking input from HDFS:

1. Create an input file in HDFS and upload it in spark as shown below

```
scala> val inputFile = sc.textFile("hdfs:///wordcountdemo/input/file1.txt")
inputFile: org.apache.spark.rdd.RDD[String] = hdfs:///wordcountdemo/input/file1.
txt MapPartitionsRDD[1] at textFile at <console>:23

scala> inputFile.foreach(println)
hello hi
Hello World of Hadoop
hello begin
end of the end

scala>
```

2. Write the Map reduce steps

```
scala> val words = inputFile.flatMap(line => line.split(" "))
words: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[2] at flatMap at <con
sole>:23

scala> val wordPairs = words.map(word => (word, 1))
wordPairs: org.apache.spark.rdd.RDD[(String, Int)] = MapPartitionsRDD[3] at map
at <console>:23

scala> val wordCounts = wordPairs.reduceByKey(_ + _)
wordCounts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[4] at reduceBy
Key at <console>:23

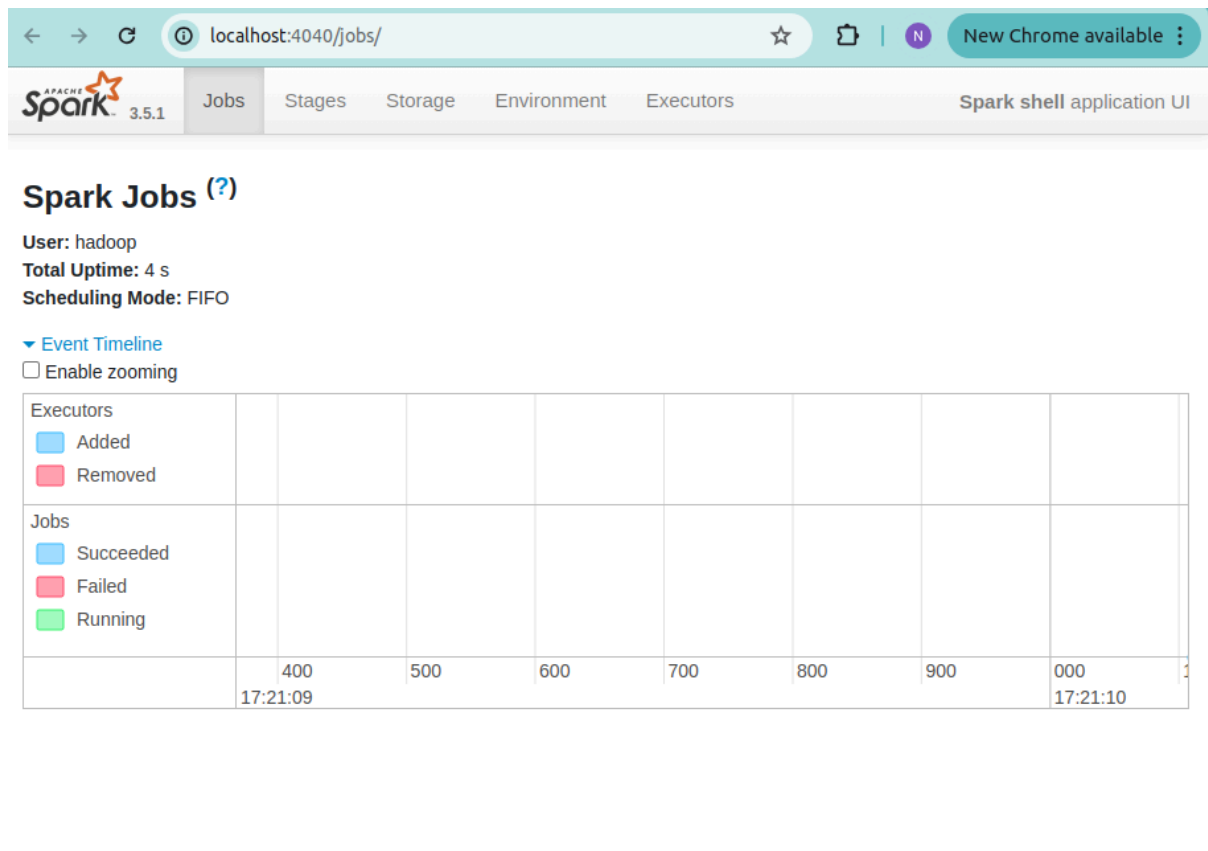
scala> wordCounts.saveAsTextFile("hdfs:///wordcountdemo/output_WC")
```

Verifying that all the files exist

```
scala> hadoop@hadoop-clone-12:~$
hadoop@hadoop-clone-12:~$ hdfs dfs -ls /wordcountdemo/output_WC
Found 3 items
-rw-r--r--  1 hadoop supergroup          0 2025-04-02 17:07 /wordcountdemo/outp
ut_WC/_SUCCESS
-rw-r--r--  1 hadoop supergroup        30 2025-04-02 17:07 /wordcountdemo/outp
ut_WC/part-00000
-rw-r--r--  1 hadoop supergroup        51 2025-04-02 17:07 /wordcountdemo/outp
ut_WC/part-00001
hadoop@hadoop-clone-12:~$
```


3. Observing the DAG visualization at

<http://localhost:4040/jobs/>



Skipped stage:

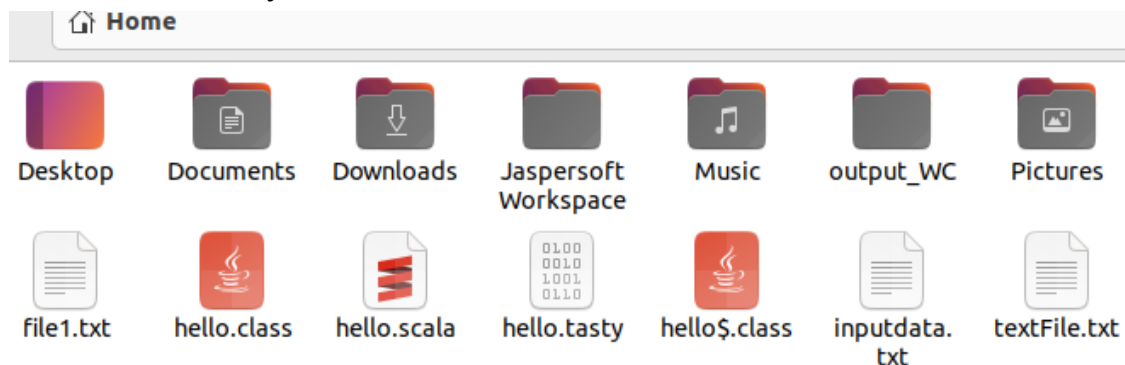
It means that data has been fetched from cache and there was no need to re-execute given stage.

It is consistent with your DAG which shows that the next stage requires shuffling (reduceByKey). Whenever there is shuffling involved Spark automatically caches generated data.

Output on hdfs looks like:

```
scala> hadoop@hadoop-clone-12:~$  
hadoop@hadoop-clone-12:~$ hdfs dfs -cat /wordcountdemo/output_WC/part-00000  
hdfs dfs -cat /wordcountdemo/output_WC/part-00001  
(Hello,1)  
(World,1)  
(hello,2)  
(begin,1)  
(hi,1)  
(end,2)  
(of,2)  
(the,1)  
(Hadoop,1)  
hadoop@hadoop-clone-12:~$
```

Files created locally:



EXERCISE :-

1. Develop simple wordcount application with apache spark(using REPL and HDFS)

Creating input in HDFS:

```
hadoop@celab138-ThinkCentre-M720t:~$ echo "spark makes big data
easy" > input.txt
hdfs dfs -mkdir -p /user/spark/input
hdfs dfs -put input.txt /user/spark/input/
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using
value of HADOOP_PREFIX.
2025-04-04 09:02:50,414 WARN util.NativeCodeLoader: Unable to l
oad native-hadoop library for your platform... using builtin-ja
va classes where applicable
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using
value of HADOOP_PREFIX.
2025-04-04 09:02:51,556 WARN util.NativeCodeLoader: Unable to l
oad native-hadoop library for your platform... using builtin-ja
va classes where applicable
```

To view the result

```
hadoop@celab138-ThinkCentre-M720t:~$ hdfs dfs -cat /user/hadoop/input/input.tx
t
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP
_PREFIX.
2025-04-04 10:16:57,259 WARN util.NativeCodeLoader: Unable to load native-hado
op library for your platform... using builtin-java classes where applicable
spark makes big data easy
```

WordCount Logic (in REPL) in spark shell, Loading input file from HDFS

```
scala> val textFile = sc.textFile("hdfs://localhost:9000/user/hadoop/input/inp
ut.txt")
textFile: org.apache.spark.rdd.RDD[String] = hdfs://localhost:9000/user/hadoop
/input/input.txt MapPartitionsRDD[49] at textFile at <console>:23

scala> val wordCounts = textFile.flatMap(_.split(" ")).map(word => (word, 1)).
reduceByKey(_ + _)
wordCounts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[52] at reduc
eByKey at <console>:23

scala> wordCounts.collect().foreach(println)
(big,1)
(data,1)
(easy,1)
(spark,1)
(makes,1)
```

Saving Output to hdfs

```
scala> wordCounts.saveAsTextFile("hdfs://localhost:9000/user/spark/output")
```

Verifying in hdfs

```
hadoop@celab138-ThinkCentre-M720t:~$ hdfs dfs -ls /user/hadoop/output_wc
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
2025-04-04 10:21:05,600 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
-rw-r--r--  3 hadoop supergroup          0 2025-04-04 10:20 /user/hadoop/output_wc/_SUCCESS
-rw-r--r--  3 hadoop supergroup        26 2025-04-04 10:20 /user/hadoop/output_wc/part-00000
-rw-r--r--  3 hadoop supergroup        20 2025-04-04 10:20 /user/hadoop/output_wc/part-00001
hadoop@celab138-ThinkCentre-M720t:~$ hdfs dfs -cat /user/hadoop/output_wc/part-00000
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
2025-04-04 10:21:13,404 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
(big,1)
(data,1)
(easy,1)
```

2. Develop Wordcount program using cluster mode(standalone or YARN mode)

Spark Standalone mode:

We can launch a standalone cluster either manually, by starting a master and workers by hand, or use our provided launch scripts. It is also possible to run these daemons on a single machine for testing.

Creating WordCount.scala at /home/hadoop

```
Open  WordCount.scala
1 import org.apache.spark.{SparkConf, SparkContext}
2
3 object WordCount {
4   def main(args: Array[String]): Unit = {
5     val conf = new SparkConf().setAppName("WordCount")
6     val sc = new SparkContext(conf)
7
8     val textFile = sc.textFile(args(0))
9     val wordCounts = textFile.flatMap(_.split(" "))
10      .map(word => (word, 1))
11      .reduceByKey(_ + _)
12
13     wordCounts.saveAsTextFile(args(1))
14     sc.stop()
15   }
16 }
```

Submitting a Scala script directly

```
hadoop@celab138-ThinkCentre-M720t:~$ spark-submit --master spark://localhost:7077 /home/hadoop/WordCount.scala
25/04/04 09:16:50 WARN Utils: Your hostname, celab138-ThinkCentre-M720t resolves to a loopback address: 127.0.1.1; using 192.168.37.16 instead (on interface eno1)
25/04/04 09:16:50 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
25/04/04 09:16:50 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Exception in thread "main" org.apache.spark.SparkException: Failed to get main class in JAR with error 'null'. Please specify one with --class.
    at org.apache.spark.deploy.SparkSubmit.error(SparkSubmit.scala:1038)
    at org.apache.spark.deploy.SparkSubmit.prepareSubmitEnvironment(SparkSubmit.scala:524)
    at org.apache.spark.deploy.SparkSubmit.org$apache$spark$deploy$SparkSubmit$$runMain(SparkSubmit.scala:955)
    at org.apache.spark.deploy.SparkSubmit.doRunMain$1(SparkSubmit.scala:192)
    at org.apache.spark.deploy.SparkSubmit.submit(SparkSubmit.scala:215)
    at org.apache.spark.deploy.SparkSubmit.doSubmit(SparkSubmit.scala:91)
    at org.apache.spark.deploy.SparkSubmit$$anon$2.doSubmit(SparkSubmit.scala:1111)
    at org.apache.spark.deploy.SparkSubmit$.main(SparkSubmit.scala:1120)
    at org.apache.spark.deploy.SparkSubmit.main(SparkSubmit.scala)
```

There are 2 ways to execute the program:

1) We can use IDE and create a jar file and then submit the jar file to the master for execution

I. Now create a jar file in IDE or CLI

II. Run the jar file and check the output status

III. Submit the application to spark so that spark can execute the application

2) use CLI to create a scala program and then run it directly on the shell

Refer the section of creating scala script

I am using way 1) here:

Compiling in jar

```
hadoop@celab138-ThinkCentre-M720t:~$ scalac -classpath "$(find $SPARK_HOME/jars -name '*.jar' | tr '\n' ':')" WordCount.scala

hadoop@celab138-ThinkCentre-M720t:~$ jar cf WordCount.jar WordCount*.class
```

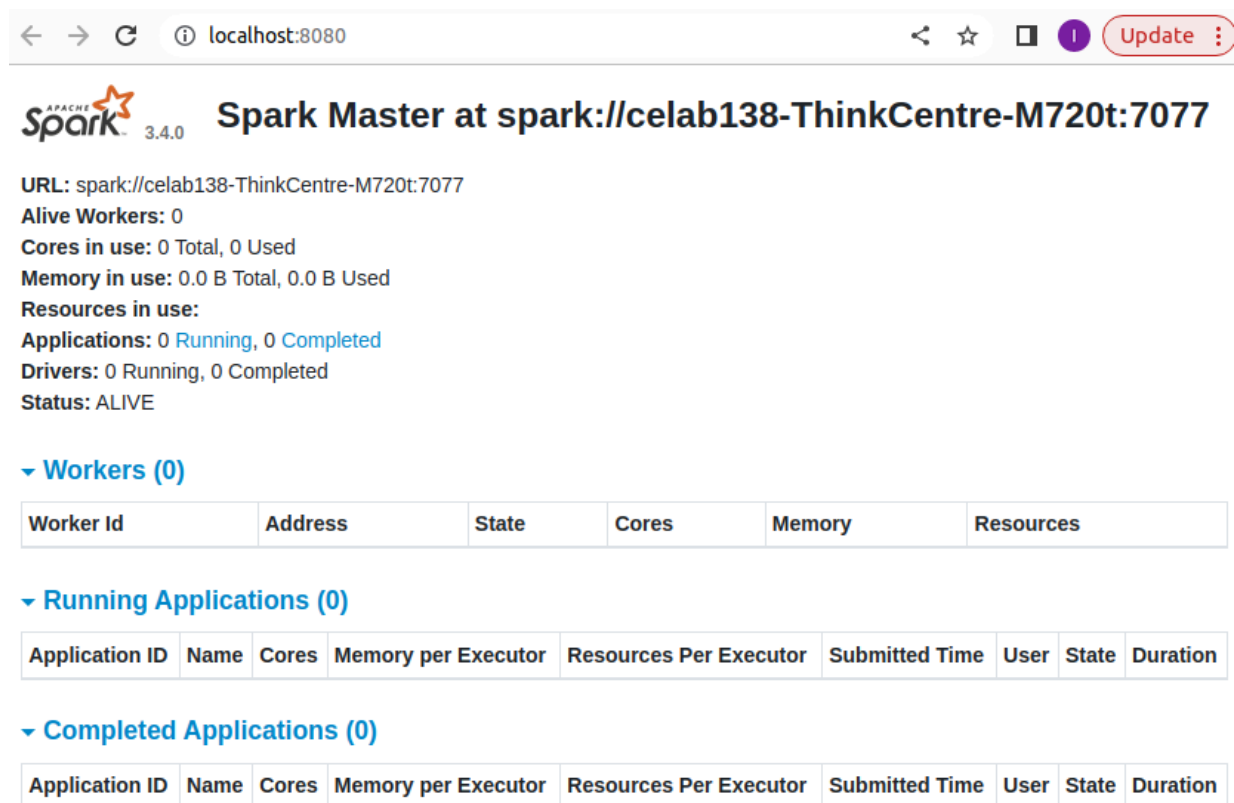
Now we start Spark Master and Spark Worker as below

```
hadoop@celab138-ThinkCentre-M720t:~$ $SPARK_HOME/sbin/start-master.sh
starting org.apache.spark.deploy.master.Master, logging to /opt/spark/logs/spark-hadoop-0
rg.apache.spark.deploy.master.Master-1-celab138-ThinkCentre-M720t.out
hadoop@celab138-ThinkCentre-M720t:~$ $SPARK_HOME/sbin/start-worker.sh spark://localhost:7077
starting org.apache.spark.deploy.worker.Worker, logging to /opt/spark/logs/spark-hadoop-0
rg.apache.spark.deploy.worker.Worker-1-celab138-ThinkCentre-M720t.out
```

Verifying with jps

```
hadoop@celab138-ThinkCentre-M720t:~$ jps
9264 Worker
5696 SecondaryNameNode
5937 ResourceManager
6067 NodeManager
5303 NameNode
6697 SparkSubmit
5449 DataNode
9339 Jps
9055 Master
```

I can now Access Spark UI on <http://localhost:8080>



← → ↻ ⓘ localhost:8080 🔗 ☆ 📄 ⓘ Update ⋮

APACHE SPARK 3.4.0 **Spark Master at spark://celab138-ThinkCentre-M720t:7077**

URL: spark://celab138-ThinkCentre-M720t:7077
Alive Workers: 0
Cores in use: 0 Total, 0 Used
Memory in use: 0.0 B Total, 0.0 B Used
Resources in use:
Applications: 0 Running, 0 Completed
Drivers: 0 Running, 0 Completed
Status: ALIVE

▼ **Workers (0)**

Worker Id	Address	State	Cores	Memory	Resources
-----------	---------	-------	-------	--------	-----------

▼ **Running Applications (0)**

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	------------------------	----------------	------	-------	----------

▼ **Completed Applications (0)**

Application ID	Name	Cores	Memory per Executor	Resources Per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	------------------------	----------------	------	-------	----------

Loading and running the scala file from shell:

Also, verifying the output on scala as below

```
scala> val textFile = sc.textFile("hdfs://localhost:9000/user/hadoop/input/input.txt")
textFile: org.apache.spark.rdd.RDD[String] = hdfs://localhost:9000/user/hadoop/input/input.txt MapPartitionsRDD[19] at textFile at <console>:23

scala> val wordCounts = textFile.flatMap(_.split(" ")).map(word => (word, 1)).reduceByKey(_ + _)
wordCounts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[22] at reduceByKey at <console>:23

scala> wordCounts.collect().foreach(println)
(is,1)
(hello,3)
(world,1)
(running,1)
(spark,2)
(hadoop,1)
```

3. Develop any custom map reduce problem using apache spark and hadoop.

We have taken a custom problem: **Character Frequency Count**

Loading the Input File from HDFS, and checking it afterwards by using:

`textFile.collect().foreach(println)`

```
scala> val textFile = sc.textFile("hdfs://localhost:9000/user/hadoop/input/input.txt")
textFile: org.apache.spark.rdd.RDD[String] = hdfs://localhost:9000/user/hadoop/input/input.txt MapPartitionsRDD[19] at textFile at <console>:23

scala> textFile.collect().foreach(println)
hello hadoop hello spark hello world
```

Checking if the file exists in HDFS

```
hadoop@celab138-ThinkCentre-M720t:~$ hdfs dfs -ls /user/hadoop/charcount/
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
2025-04-04 09:38:28,917 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
-rw-r--r--  3 hadoop supergroup          0 2025-04-04 09:38 /user/hadoop/charcount/_SUCCESS
-rw-r--r--  3 hadoop supergroup        37 2025-04-04 09:38 /user/hadoop/charcount/part-00000
-rw-r--r--  3 hadoop supergroup          0 2025-04-04 09:38 /user/hadoop/charcount/part-00001
```


Printing output on hdfs

```
hadoop@celab138-ThinkCentre-M720t:~$ hdfs dfs -cat /user/hadoop/charcount/part-00000
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX
2025-04-04 09:38:35,978 WARN util.NativeCodeLoader: Unable to load native-hadoop library
hello hadoop hello spark hello world
```

FlatMap to Characters (transforming the text into individual characters) :

```
scala> val characters = textFile.flatMap(line => line.replaceAll(" ", "").toCharArray)
characters: org.apache.spark.rdd.RDD[Char] = MapPartitionsRDD[20] at flatMap at <console>:23

scala> characters.collect().foreach(println)
h
e
l
l
o
h
a
d
o
o
p
h
e
l
l
o
s
p
a
r
k
h
e
l
l
o
w
o
r
l
d
```


Converting to Key-Value Pairs and checking the output:

```
scala> val charPairs = characters.map(char => (char, 1))
charPairs: org.apache.spark.rdd.RDD[(Char, Int)] = MapPartitionsRDD[21] at map
at <console>:23

scala> charPairs.collect().foreach(println)
(h,1)
(e,1)
(l,1)
(l,1)
(o,1)
(h,1)
(a,1)
(d,1)
(o,1)
(o,1)
(p,1)
(h,1)
(e,1)
(l,1)
(l,1)
(o,1)
(s,1)
(p,1)
(a,1)
(r,1)
(k,1)
(h,1)
(e,1)
(l,1)
(l,1)
(o,1)
(w,1)
(o,1)
(r,1)
(l,1)
(d,1)
```

Performing ReduceByKey to Count

```
scala> val charCounts = charPairs.reduceByKey(_ + _)
charCounts: org.apache.spark.rdd.RDD[(Char, Int)] = ShuffledRDD[22] at reduceB
yKey at <console>:23

scala> charCounts.collect().foreach(println)
(d,2)
(p,2)
(h,4)
(r,2)
(l,7)
(w,1)
(s,1)
(e,3)
(a,2)
(k,1)
(o,6)
```

Saving Output to HDFS

```
scala> charCounts.saveAsTextFile("hdfs://localhost:9000/user/hadoop/charcount_fixed")
```

Verifying saved data:

```
hadoop@celab138-ThinkCentre-M720t:~$ hdfs dfs -ls /user/hadoop/charcount_fixed
/
hdfs dfs -cat /user/hadoop/charcount_fixed/part-00000
hdfs dfs -cat /user/hadoop/charcount_fixed/part-00001
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
2025-04-04 09:46:58,555 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
-rw-r--r--    3 hadoop supergroup          0 2025-04-04 09:46 /user/hadoop/charcount_fixed/_SUCCESS
-rw-r--r--    3 hadoop supergroup        30 2025-04-04 09:46 /user/hadoop/charcount_fixed/part-00000
-rw-r--r--    3 hadoop supergroup        36 2025-04-04 09:46 /user/hadoop/charcount_fixed/part-00001
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
2025-04-04 09:46:59,618 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
(d,2)
(p,2)
(h,4)
(r,2)
(l,7)
WARNING: HADOOP_PREFIX has been replaced by HADOOP_HOME. Using value of HADOOP_PREFIX.
2025-04-04 09:47:00,774 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
(w,1)
(s,1)
(e,3)
(a,2)
(k,1)
(o,6)
```

To visualize DAG

```
scala> charCounts.toDebugString
res17: String =
(2) ShuffledRDD[22] at reduceByKey at <console>:23 []
+- (2) MapPartitionsRDD[21] at map at <console>:23 []
    | MapPartitionsRDD[20] at flatMap at <console>:23 []
    | hdfs://localhost:9000/user/hadoop/input/input.txt MapPartitionsRDD[19]
at textFile at <console>:23 []
    | hdfs://localhost:9000/user/hadoop/input/input.txt HadoopRDD[18] at text
File at <console>:23 []
```

Spark UI is present at <http://localhost:4040/jobs/>

Spark Jobs (?)

User: hadoop
 Total Uptime: 12 min
 Scheduling Mode: FIFO
 Completed Jobs: 11

▶ Event Timeline

▼ Completed Jobs (11)

Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

Job Id ▼	Description	Submitted	Duration	Stages: Succeeded/Total	Tasks (for all stages): Succeeded/Total
10	runJob at SparkHadoopWriter.scala:83 runJob at SparkHadoopWriter.scala:83	2025/04/04 09:46:05	0.2 s	1/1 (1 skipped)	2/2 (2 skipped)
9	collect at <console>:24 collect at <console>:24	2025/04/04 09:45:56	42 ms	2/2	4/4
8	collect at <console>:24 collect at <console>:24	2025/04/04 09:45:47	16 ms	1/1	2/2
7	collect at <console>:24 collect at <console>:24	2025/04/04 09:45:31	27 ms	1/1	2/2
6	collect at <console>:24 collect at <console>:24	2025/04/04 09:45:10	25 ms	1/1	2/2
5	collect at <console>:24 collect at <console>:24	2025/04/04 09:44:50	16 ms	1/1	2/2
4	sortBy at <console>:24 sortBy at <console>:24	2025/04/04 09:44:46	58 ms	2/2	4/4

localhost:4040/jobs/job/?id=10

Update

Details for Job 10

Status: SUCCEEDED

Submitted: 2025/04/04 09:46:05

Duration: 0.2 s

Completed Stages: 1

Skipped Stages: 1

▶ Event Timeline

▼ DAG Visualization

Stage 13 (skipped)

textFile

flatMap

map

Stage 14

reduceByKey

saveAsTextFile

▼ Completed Stages (1)

Page: 1

1 Pages. Jump to 1. Show 100 items in a page.

Stage Id ▼	Description	Submitted	Duration	Tasks: Succeeded/Total	Input	Output	Shuffle Read	Shuffle Write
14	runJob at SparkHadoopWriter.scala:83	2025/04/04 09:46:05	0.1 s	2/2		66.0 B	161.0 B	

Big Data Analytics

Details for Stage 14 (Attempt 0)

Resource Profile Id: 0

Total Time Across All Tasks: 0.2 s

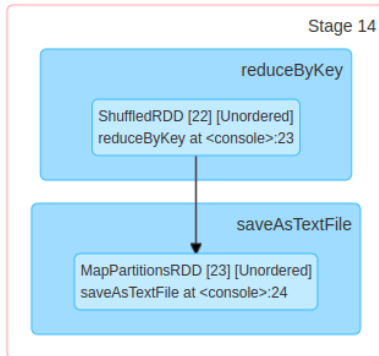
Locality Level Summary: Node local: 2

Output Size / Records: 66.0 B / 11

Shuffle Read Size / Records: 161.0 B / 11

Associated Job Ids: 10

▼ DAG Visualization



► Show Additional Metrics

► Event Timeline

Summary Metrics for 2 Completed Tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	0.1 s	0.1 s	0.1 s	0.1 s	0.1 s
GC Time	0.0 ms	0.0 ms	0.0 ms	0.0 ms	0.0 ms
Output Size / Records	30 B / 5	30 B / 5	36 B / 6	36 B / 6	36 B / 6
Shuffle Read Size / Records	77 B / 5	77 B / 5	84 B / 6	84 B / 6	84 B / 6





Also, The file system on <http://localhost:50070> looks like:

localhost:50070/explorer.html#/




Hadoop

Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Browse Directory

/ Go!    

Show 25 ▾ entries Search:

<input type="checkbox"/>	↕	Permission	↕	Owner	↕	Group	↕	Size	↕	Last Modified	↕	Replication	↕	Block Size	↕	Name	↕
<input type="checkbox"/>		drwxr-xr-x		hadoop		supergroup		0 B		Oct 03 2023		0		0 B		spark-tut	
<input type="checkbox"/>		drwxrwx---		hadoop		supergroup		0 B		May 10 2023		0		0 B		tmp	
<input type="checkbox"/>		drwxr-xr-x		hadoop		supergroup		0 B		Apr 04 09:30		0		0 B		user	

Showing 1 to 3 of 3 entries

Previous 1 Next

In spark-tut:

→ ↻ ⓘ localhost:50070/explorer.html#/spark-tut 🔍 ☆ 📄 ⓘ Update

Hadoop

Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Browse Directory

/spark-tut Go! 📁 ⬆️ 📄 🗑️

Show 25 ▾ entries Search:

<input type="checkbox"/>	↕	Permission	↕	Owner	↕	Group	↕	Size	↕	Last Modified	↕	Replication	↕	Block Size	↕	Name	↕
<input type="checkbox"/>		-rw-r--r--		hadoop		supergroup		25 B		Oct 03 2023		3		128 MB		inputdata.txt	🗑️

Showing 1 to 1 of 1 entries

Previous 1 Next

Hadoop, 2023.

In user:

→ ↻ ⓘ localhost:50070/explorer.html#/user 🔍 ☆ 📄 ⓘ Update

Hadoop

Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Browse Directory

/user Go! 📁 ⬆️ 📄 🗑️

Show 25 ▾ entries Search:

<input type="checkbox"/>	↕	Permission	↕	Owner	↕	Group	↕	Size	↕	Last Modified	↕	Replication	↕	Block Size	↕	Name	↕
<input type="checkbox"/>		drwxr-xr-x		hadoop		supergroup		0 B		Apr 04 09:46		0		0 B		hadoop	🗑️
<input type="checkbox"/>		drwxr-xr-x		hadoop		supergroup		0 B		Apr 04 09:14		0		0 B		spark	🗑️

Showing 1 to 2 of 2 entries

Previous 1 Next





hadoop:

→ ↻ ⓘ localhost:50070/explorer.html#/user/hadoop < ☆ □ ⓘ Update ⋮













Hadoop

Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Browse Directory

/user/hadoop Go!    

Show 25 ▾ entries Search:

<input type="checkbox"/>	 Permission	 Owner	 Group	 Size	 Last Modified	 Replication	 Block Size	 Name	
<input type="checkbox"/>	drwxr-xr-x	hadoop	supergroup	0 B	Apr 04 09:38	0	0 B	charcount	
<input type="checkbox"/>	drwxr-xr-x	hadoop	supergroup	0 B	Apr 04 09:46	0	0 B	charcount_fixed	
<input type="checkbox"/>	drwxr-xr-x	hadoop	supergroup	0 B	Apr 04 09:31	0	0 B	input	

Showing 1 to 3 of 3 entries





Previous 1 Next

→ ↻ localhost:50070/explorer.html#/user/hadoop/charcount Update




Hadoop

Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Browse Directory

/user/hadoop/charcount Go!    

Show 25 entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	-rw-r--r--	hadoop	supergroup	0 B	Apr 04 09:38	3	128 MB	_SUCCESS	
<input type="checkbox"/>	-rw-r--r--	hadoop	supergroup	37 B	Apr 04 09:38	3	128 MB	part-00000	
<input type="checkbox"/>	-rw-r--r--	hadoop	supergroup	0 B	Apr 04 09:38	3	128 MB	part-00001	

Showing 1 to 3 of 3 entries Previous 1 Next

[Download](#) [Head the file \(first 32K\)](#) [Tail the file \(last 32K\)](#)

Block information -- Block 0

Block ID: 1073741837
Block Pool ID: BP-1266388418-127.0.0.1-1560101340487
Generation Stamp: 1013
Size: 30
Availability:
• localhost

File contents

```
(Hello,1)
(World,1)
(hello,2)
```

spark:

→ ↻ ⓘ localhost:50070/explorer.html#/user/spark ← ☆ 🗄 ⓘ Update ⋮

Hadoop

Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Browse Directory

/user/spark Go! 📁 🔍 📄 📂

Show 25 ▾ entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	drwxr-xr-x	hadoop	supergroup	0 B	Apr 04 09:02	0	0 B	input	🗑
<input type="checkbox"/>	drwxr-xr-x	hadoop	supergroup	0 B	Apr 04 09:07	0	0 B	output	🗑
<input type="checkbox"/>	drwxr-xr-x	hadoop	supergroup	0 B	Apr 04 09:14	0	0 B	output-wordcount	🗑

Showing 1 to 3 of 3 entries Previous **1** Next

→ ↺

localhost:50070/explorer.html#/user/spark/output

← ☆ □ ⓘ Update

Hadoop

Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities ▾

Browse Directory

/user/spark/output

Go!

📁

🔍

📄

🗑️

Show

25 ▾

 entries Search:

<input type="checkbox"/>	↕	Permission	↕	Owner	↕	Group	↕	Size	↕	Last Modified	↕	Replication	↕	Block Size	↕	Name	↕
<input type="checkbox"/>		-rw-r--r--		hadoop		supergroup		0 B		Apr 04 09:07		3		128 MB		_SUCCESS	🗑️
<input type="checkbox"/>		-rw-r--r--		hadoop		supergroup		26 B		Apr 04 09:07		3		128 MB		part-00000	🗑️
<input type="checkbox"/>		-rw-r--r--		hadoop		supergroup		0 B		Apr 04 09:07		3		128 MB		part-00001	🗑️

Showing 1 to 3 of 3 entries

Previous

1

 Next

Download

Head the file (first 32K)

Tail the file (last 32K)

Block information ▾

Block 0 ▾

Block ID: 1073741838
Block Pool ID: BP-1266388418-127.0.0.1-1560101340487
Generation Stamp: 1014
Size: 59
Availability:

- localhost

File contents

```
(begin,2)
(hi,1)
(end,2)
(of,2)
(enf,1)
(the,1)
(Hadoop,1)
```

Summarised learning:

In this lab, I explored the architecture and execution model of Apache Spark for big data analytics. I developed and executed a WordCount application using both REPL and HDFS, and ran it in cluster mode using standalone Spark deployment. I also created a custom MapReduce program to perform character frequency analysis, showcasing Spark's flexibility for scalable data transformations. By integrating HDFS as the data source, I implicitly understood the coordination between Spark and Hadoop components. This hands-on experience deepened my understanding of Spark's multi-language support, distributed processing, and practical use for real-world data analytics tasks.