

LAB - 05

Aim:- To learn and understand the use of mongodb to create and analyze the data.

Connect to a Local Deployment on the Default Port:

mongosh

To show list of data bases available in the system use following command:

> show dbs

Create a New Database:

> use product:

To display the name of current working Database:

> db

```
hadoop@hadoop-clone-11:~$ mongosh
Current Mongosh Log ID: 67d163569fd0e214bca26a12
Connecting to:      mongodb://127.0.0.1:27017/?directConne
Using MongoDB:      7.0.9
Using Mongosh:      2.2.6
mongosh 2.4.2 is available for download: https://www.mongodb.com

For mongosh info see: https://docs.mongodb.com/mongodb-shell/

-----
The server generated these startup warnings when booting
2025-03-12T15:39:08.149+05:30: Using the XFS filesystem is
2025-03-12T15:39:08.759+05:30: Access control is not enable
-----

test> show dbs
admin      40.00 KiB
config    108.00 KiB
local      96.00 KiB
test> db
test
test> use product
switched to db product
```

To Create a Collection

To Display a Collections

show collections

```
product> show collections

product> db.books.insertMany([
...   {name: 'The Complete Reference', page: 500, price: 1000},
...   {name: 'The Python Programming', page: 800, price: 1500},
...   {name: 'Software Engineering', page: 500, price: 2000}
... ])
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId('67d163d29fd0e214bca26a13'),
    '1': ObjectId('67d163d29fd0e214bca26a14'),
    '2': ObjectId('67d163d29fd0e214bca26a15')
  }
}
```

To retrieve collection

1. db.collection.find(): >db.books.find()
2. db.collection.find().pretty(): To pretify the output
>db.books.find().pretty()

```
product> db.books.find().pretty()
[
  {
    _id: ObjectId('67d163d29fd0e214bca26a13'),
    name: 'The Complete Reference',
    page: 500,
    price: 1000
  },
  {
    _id: ObjectId('67d163d29fd0e214bca26a14'),
    name: 'The Python Programming',
    page: 800,
    price: 1500
  },
  {
    _id: ObjectId('67d163d29fd0e214bca26a15'),
    name: 'Software Engineering',
    page: 500,
    price: 2000
  }
]
```

Insert Documents(insert command):

```
product> db.books.insertMany([
...   {name: 'Clean Code', page: 464, price: 1200},
...   {name: 'Introduction to Algorithms', page: 1312, price: 2500},
...   {name: 'Design Patterns', page: 395, price: 1800},
...   {name: 'The Pragmatic Programmer', page: 352, price: 1500},
...   {name: 'Artificial Intelligence: A Modern Approach', page: 1136, price: 3000}
... ])
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId('67d164109fd0e214bca26a16'),
    '1': ObjectId('67d164109fd0e214bca26a17'),
    '2': ObjectId('67d164109fd0e214bca26a18'),
    '3': ObjectId('67d164109fd0e214bca26a19'),
    '4': ObjectId('67d164109fd0e214bca26a1a')
  }
}
product> db.books.find().pretty()
```

To retrieve this collection:

```
product> db.books.find().pretty()
[
  {
    _id: ObjectId('67d163d29fd0e214bca26a13'),
    name: 'The Complete Reference',
    page: 500,
    price: 1000
  },
  {
    _id: ObjectId('67d163d29fd0e214bca26a14'),
    name: 'The Python Programming',
    page: 800,
    price: 1500
  },
  {
    _id: ObjectId('67d163d29fd0e214bca26a15'),
    name: 'Software Engineering',
    page: 500,
    price: 2000
  },
  {
    _id: ObjectId('67d164109fd0e214bca26a16'),
    name: 'Clean Code',
    page: 464,
    price: 1200
  },
  {
    _id: ObjectId('67d164109fd0e214bca26a17'),
    name: 'Introduction to Algorithms',
    page: 1312,
    price: 2500
  },
  {
    _id: ObjectId('67d164109fd0e214bca26a18'),
    name: 'Design Patterns',
    page: 395,
    price: 1800
  },
  {
    _id: ObjectId('67d164109fd0e214bca26a19'),
    name: 'The Pragmatic Programmer',
    page: 352,
    price: 1500
  },
  {
    _id: ObjectId('67d164109fd0e214bca26a1a'),
    name: 'Artificial Intelligence: A Modern Approach',
    page: 1136,
    price: 3000
  }
]
```

The database output looks like this:

localhost:27017 ... {} My Queries product books x +

localhost:27017 > product > books

Documents 3 Aggregations Schema Indexes 1 Va

Type a query: { field: 'value' } or [Generate query](#)

ADD DATA EXPORT DATA UPDATE DELETE

```

_id: ObjectId('67d1666e9fd0e214bca26a23')
name: "The Complete Reference"
page: 500
price: 1000

_id: ObjectId('67d1666e9fd0e214bca26a24')
name: "The Python Programming"
page: 800
price: 1500

_id: ObjectId('67d1666e9fd0e214bca26a25')
name: "Software Engineering"
page: 500
price: 2000

_id: ObjectId('67d166779fd0e214bca26a26')
name: "Clean Code"
page: 464
price: 1200

_id: ObjectId('67d166779fd0e214bca26a27')
name: "Introduction to Algorithms"
page: 1312
price: 2500

```

```

_id: ObjectId('67d166779fd0e214bca26a28')
name: "Design Patterns"
page: 395
price: 1800

_id: ObjectId('67d166779fd0e214bca26a29')
name: "The Pragmatic Programmer"
page: 352
price: 1500

_id: ObjectId('67d166779fd0e214bca26a2a')
name: "Artificial Intelligence: A Modern Approach"
page: 1136
price: 3000

```

Update a Single Document

```
product> db.books.updateOne(
...   { "_id" : ObjectId("64199568d7fb909707ec850b") },
...   { $set: { name: "The Complete Reference", pages: 8000, price: 5000 } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 0,
  modifiedCount: 0,
  upsertedCount: 0
}
```

Update Multiple Documents

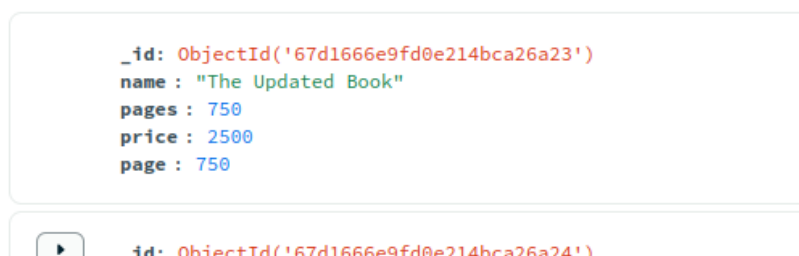

```
product> db.books.updateMany(
...   { name: "The Complete Reference" },
...   { $set: { pages: 8000, price: 5000 } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```

```
_id: ObjectId('67d1666e9fd0e214bca26a23')
name : "The Complete Reference"
page : 500
price : 5000
pages : 8000
```

Replace a Document

```
product> db.books.replaceOne(
...   { "_id" : ObjectId("67d1666e9fd0e214bca26a23") },
...   { name: "The Updated Book", pages: 750, price: 2500, page: 750 }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```

```
product> db.books.find().pretty()
[
  {
    _id: ObjectId('67d1666e9fd0e214bca26a23'),
    name: 'The Updated Book',
    pages: 750,
    price: 2500,
    page: 750
  },
  {
    _id: ObjectId('67d1666e9fd0e214bca26a24'),
    name: 'The Python Programming',
    page: 800,
    price: 1500
  },
  {
    _id: ObjectId('67d1666e9fd0e214bca26a25'),
    name: 'Software Engineering',
    page: 500,
    price: 2000
  },
  {
    _id: ObjectId('67d166779fd0e214bca26a26'),
    name: 'Clean Code',
    page: 464,
    price: 1200
  },
  {
    _id: ObjectId('67d166779fd0e214bca26a27'),
    name: 'Introduction to Algorithms'
  }
]
```




Update Operators

Increase Price (**\$inc**), Rename Field (**\$rename**), Set a New Field (**\$set**), Remove a Field (**\$unset**)

```
product> db.books.updateOne(
...   { name: "The Python Programming" },
...   { $inc: { price: 500 } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
product> db.books.updateMany(
...   {},
...   { $rename: { "page": "total_pages" } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 8,
  modifiedCount: 8,
  upsertedCount: 0
}
product> db.books.updateMany(
...   {},
...   { $set: { available: true } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 8,
  modifiedCount: 8,
  upsertedCount: 0
}
product> db.books.updateOne(
...   { name: "The Python Programming" },
...   { $unset: { price: "" } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```


The screenshot shows the MongoDB Compass web interface. The left sidebar displays the database structure: **localhost:27017** > **product** > **books**. The main panel shows the 'Documents' tab with 3 documents. A search bar and buttons for 'ADD DATA', 'EXPORT DATA', and 'UPDATE' are visible. The documents are listed as follows:

Document	_id	name	pages	price	total_pages	available
1	ObjectId('67d1666e9fd0e214bca26a23')	"The Updated Book"	750	2500	750	true
2	ObjectId('67d1666e9fd0e214bca26a24')	"The Python Programming"			800	true
3	ObjectId('67d1666e9fd0e214bca26a25')	"Software Engineering"		2000	500	true
4	ObjectId('67d166779fd0e214bca26a26')	"Clean Code"		1200	464	true

	<pre>_id: ObjectId('67d166779fd0e214bca26a27') name: "Introduction to Algorithms" price: 2500 total_pages: 1312 available: true</pre>
	<pre>_id: ObjectId('67d166779fd0e214bca26a28') name: "Design Patterns" price: 1800 total_pages: 395 available: true</pre>
	<pre>_id: ObjectId('67d166779fd0e214bca26a29') name: "The Pragmatic Programmer" price: 1500 total_pages: 352 available: true</pre>
	<pre>_id: ObjectId('67d166779fd0e214bca26a2a') name: "Artificial Intelligence: A Modern Approach" price: 3000 total_pages: 1136 available: true</pre>

Updating Documents

Update the Document with a Specific ID,

Update the Name to **bookname** for Specific ID,

Update the Name to **book_name** Where Pages = 8000,,

Update All Documents with Price = 800 to Rename Name to **book_name**,

Increase Price by 5 Where Price = 800

```
product> db.books.updateOne(
...   { _id: ObjectId("64199568d7fb909707ec850b") },
...   { $set: { name: "the complete reference 2", pages: 8000, price: 5000 } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 0,
  modifiedCount: 0,
  upsertedCount: 0
}
product> db.books.updateOne(
...   { _id: ObjectId("64199568d7fb909707ec850b") },
...   { $rename: { name: "bookname" } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 0,
  modifiedCount: 0,
  upsertedCount: 0
}
product> db.books.update(
...   { pages: 8000 },
...   { $rename: { name: "book_name" } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 0,
  modifiedCount: 0,
  upsertedCount: 0
}
product> db.books.updateMany(
...   { price: 800 },
...   { $rename: { name: "book_name" } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 0,
  modifiedCount: 0,
  upsertedCount: 0
}
product> db.books.update(
...   { price: 800 },
...   { $inc: { price: 5 } }
... )
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 0,
  modifiedCount: 0,
  upsertedCount: 0
}
```

Deleting Documents

Delete Multiple Documents Where Title = Titanic,

Remove Document by ID,

Delete One Document Where Cast = Brad Pitt

```
product> db.movies.deleteMany({ title: "Titanic" })
{ acknowledged: true, deletedCount: 0 }
product> db.books.remove({ _id: ObjectId("64199568d7fb909707ec850b") })
DeprecationWarning: Collection.remove() is deprecated. Use deleteOne, deleteMany,
  findOneAndDelete, or bulkWrite.
{ acknowledged: true, deletedCount: 0 }
product> db.movies.deleteOne({ cast: "Brad Pitt" })
{ acknowledged: true, deletedCount: 0 }
product> db.books.find()
[
  {
    _id: ObjectId('67d1666e9fd0e214bca26a23'),
    name: 'The Updated Book',
    pages: 750,
    price: 2500,
    total_pages: 750,
    available: true
  },
  {
    _id: ObjectId('67d1666e9fd0e214bca26a24'),
    name: 'The Python Programming',
    total_pages: 800,
    available: true
  }
]
```

Querying Documents

```
product> db.books.find().limit(2)
[
  {
    _id: ObjectId('67d1666e9fd0e214bca26a23'),
    name: 'The Updated Book',
    pages: 750,
    price: 2500,
    total_pages: 750,
    available: true
  },
  {
    _id: ObjectId('67d1666e9fd0e214bca26a24'),
    name: 'The Python Programming',
    total_pages: 800,
    available: true
  }
]
product> db.books.find().sort({ price: 1 }).pretty()
[
  {
    _id: ObjectId('67d1666e9fd0e214bca26a24'),
    name: 'The Python Programming',
    total_pages: 800,
    available: true
  },
  {
    _id: ObjectId('67d166779fd0e214bca26a26'),
    name: 'Clean Code',
    price: 1200
  }
]
```

```
product> db.books.find().limit(3).sort({ name: -1 }).pretty()
[
  {
    _id: ObjectId('67d1666e9fd0e214bca26a23'),
    name: 'The Updated Book',
    pages: 750,
    price: 2500,
    total_pages: 750,
    available: true
  },
  {
    _id: ObjectId('67d1666e9fd0e214bca26a24'),
    name: 'The Python Programming',
    total_pages: 800,
    available: true
  },
  {
    _id: ObjectId('67d166779fd0e214bca26a29'),
    name: 'The Pragmatic Programmer',
    price: 1500,
    total_pages: 352,
    available: true
  }
]
product> db.books.find({ price: { $eq: 800 } })

product> db.books.find({ price: { $gt: 800 } })
[
  {
    _id: ObjectId('67d1666e9fd0e214bca26a23'),
    name: 'The Updated Book',
```

```
product> db.books.find({ price: { $in: [800, 900] } })

product> db.books.find({ $and: [{ page: 600 }, { price: 800 } ] })

product> db.books.find({ price: { $not: { $gt: 800 } } })
[
  {
    _id: ObjectId('67d1666e9fd0e214bca26a24'),
    name: 'The Python Programming',
    total_pages: 800,
    available: true
  }
]
```

Projection

Projection is used to project the number of columns or attributes from a collection

Get The total number of collections

```
db.books.find().count()
```

```
product> db.books.find({}, { book_name: 1 })
[
  { _id: ObjectId('67d1666e9fd0e214bca26a23') },
  { _id: ObjectId('67d1666e9fd0e214bca26a24') },
  { _id: ObjectId('67d1666e9fd0e214bca26a25') },
  { _id: ObjectId('67d166779fd0e214bca26a26') },
  { _id: ObjectId('67d166779fd0e214bca26a27') },
  { _id: ObjectId('67d166779fd0e214bca26a28') },
  { _id: ObjectId('67d166779fd0e214bca26a29') },
  { _id: ObjectId('67d166779fd0e214bca26a2a') }
]
product> db.books.find({ price: 805 }, { book_name: 1 })

product> db.books.find({ price: 805 }, { book_name: 1, _id: 0 })

product> db.books.find({ price: 1800 }, { Design Patterns: 1, _id: 0 })
Uncaught:
SyntaxError: Unexpected token, expected "," (1:40)

> 1 | db.books.find({ price: 1800 }, { Design Patterns: 1, _id: 0 })
    |                                     ^
    2 |

product> db.books.find({ price: 1800, name: "Design Patterns" }, { name: 1, _id:
0 })
[ { name: 'Design Patterns' } ]
product> db.books.find().count()
8
```

MongoDb with mapreduce

Consider a stud collection with following data

```
product> db.stud.mapReduce(
...   function() { emit(this.sec, this.marks); },
...   function(key, values) { return Array.sum(values); },
...   { out: "total_marks_by_section" }
... )
DeprecationWarning: Collection.mapReduce() is deprecated. Use an aggregation instead.
See https://docs.mongodb.com/manual/core/map-reduce for details.
{ result: 'total_marks_by_section', ok: 1 }
product> db.stud.find({})

product> use stud
switched to db stud
stud> show collections

stud> db.stud.insertMany([
...   { "id": 1, "sec": "A", "marks": 80 },
...   { "id": 2, "sec": "A", "marks": 90 },
...   { "id": 1, "sec": "B", "marks": 99 },
...   { "id": 1, "sec": "B", "marks": 95 },
...   { "id": 1, "sec": "C", "marks": 90 }
... ]);
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId('67d16ea39fd0e214bca26a2b'),
    '1': ObjectId('67d16ea39fd0e214bca26a2c'),
    '2': ObjectId('67d16ea39fd0e214bca26a2d'),
    '3': ObjectId('67d16ea39fd0e214bca26a2e'),
    '4': ObjectId('67d16ea39fd0e214bca26a2f')
  }
}
stud> use product
switched to db product
product> db.stud.insertMany([
...   { "id": 1, "sec": "A", "marks": 80 },
...   { "id": 2, "sec": "A", "marks": 90 },
...   { "id": 1, "sec": "B", "marks": 99 },
...   { "id": 1, "sec": "B", "marks": 95 },
...   { "id": 1, "sec": "C", "marks": 90 }
... ]);
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId('67d16eb99fd0e214bca26a30'),
    '1': ObjectId('67d16eb99fd0e214bca26a31'),
```



```
product> show collections
books
stud
total_marks_by_section
product> db.stud.find().pretty()
[
  {
    _id: ObjectId('67d16eb99fd0e214bca26a30'),
    id: 1,
    sec: 'A',
    marks: 80
  },
  {
    _id: ObjectId('67d16eb99fd0e214bca26a31'),
    id: 2,
    sec: 'A',
    marks: 90
  },
  {
    _id: ObjectId('67d16eb99fd0e214bca26a32'),
    id: 1,
    sec: 'B',
    marks: 99
  },
  {
    _id: ObjectId('67d16eb99fd0e214bca26a33'),
    id: 1,
    sec: 'B',
    marks: 95
  },
  {
    _id: ObjectId('67d16eb99fd0e214bca26a34'),
    id: 1,
    sec: 'C',
    marks: 90
  }
]
```

1. Define the map function
2. Define the reduce function
3. Use the mapReduce function by passing mapper function, reducer function and the output collection name

```
product> // 1. Define map function

product> var mapFunction = function() {
...     emit(this.id, this.marks);
... };

product> // 2. Define reduce function

product> var reduceFunction = function(key, values) {
...     return Array.sum(values);
... }; // 3. Run mapReduce

product> db.stud.mapReduce(
...     mapFunction,
...     reduceFunction,
...     { out: "student_marks_total" }
... );
{ result: 'student_marks_total', ok: 1 }
product> // Check the output

product> db.student_marks_total.find().pretty();
[ { _id: 1, value: 364 }, { _id: 2, value: 90 } ]
product> var mapFunction = function() {
...     emit(this.sec, this.marks);
... };

product> var reduceFunction = function(key, values) {
...     return Array.sum(values); // Summing up all the marks in each section
... };

product> db.stud.mapReduce(
...     mapFunction,
...     reduceFunction,
...     { out: "section_marks_summary" }
... );
{ result: 'section_marks_summary', ok: 1 }
```

```
product> db.student_marks_total.find().pretty();
[ { _id: 1, value: 364 }, { _id: 2, value: 90 } ]
product> var mapFunction = function() {
...     emit(this.sec, this.marks);
... };

product> var reduceFunction = function(key, values) {
...     return Array.sum(values); // Summing up all the marks in each section
... };

product> db.stud.mapReduce(
...     mapFunction,
...     reduceFunction,
...     { out: "section_marks_summary" }
... );
{ result: 'section_marks_summary', ok: 1 }
product> db.section_marks_summary.find().forEach(printjson);
{
  _id: 'C',
  value: 90
}
{
  _id: 'A',
  value: 170
}
{
  _id: 'B',
  value: 194
}
```



The screenshot shows the MongoDB Compass interface. On the left, the 'Databases' sidebar lists 'admin', 'config', 'local', 'product', and 'section_marks_sum...'. The 'product' database is selected. The main area shows the 'Documents' tab with 3 documents. The documents are:


- `{ "_id": "C", "value": 90 }`
- `{ "_id": "A", "value": 170 }`
- `{ "_id": "B", "value": 194 }`


The interface includes a search bar, a query input field with a placeholder 'Type a query: { field: 'value' } or [Generate query](#)', and buttons for 'ADD DATA', 'EXPORT DATA', 'UPDATE', and 'DELETE'.


My Queries


Performance


Databases  


 admin


 config



 local


 startup_log


 product

 books

 section_marks_summary

 **stud** 

 student_marks_total

 total_marks_by_section


Documents 5


Aggregations


Schema


Indexes 1

Type a query: { field: 'value' } or [Generate query](#)

 ADD DATA

 EXPORT DATA

 UPDATE

 DELETE

_id: ObjectId('67d16eb99fd0e214bca26a30')

id: 1

sec: "A"

marks: 80

_id: ObjectId('67d16eb99fd0e214bca26a31')

id: 2

sec: "A"

marks: 90

_id: ObjectId('67d16eb99fd0e214bca26a32')

id: 1

sec: "B"

marks: 99

_id: ObjectId('67d16eb99fd0e214bca26a33')

id: 1

sec: "B"

marks: 95

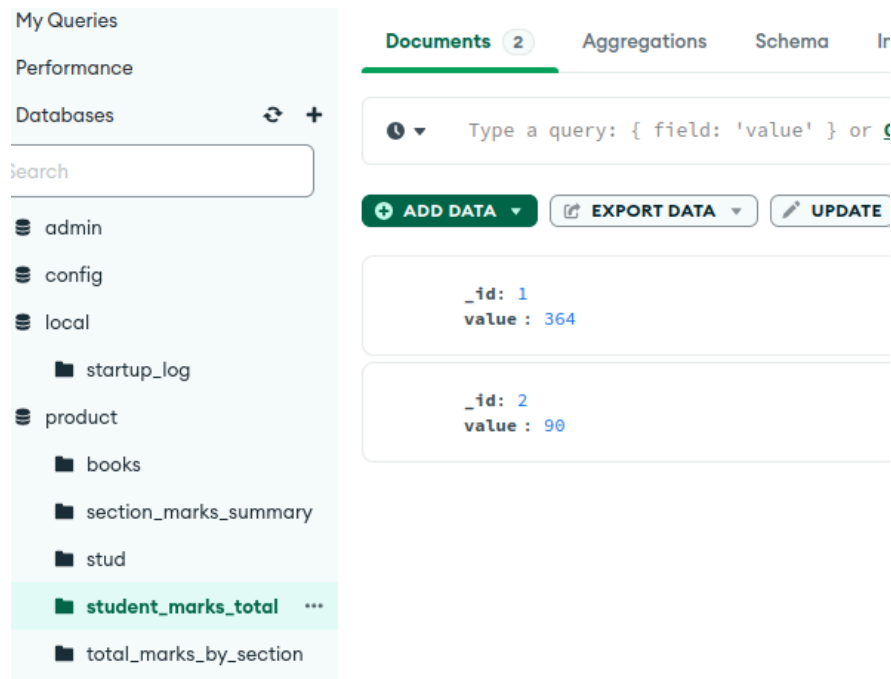
_id: ObjectId('67d16eb99fd0e214bca26a34')

id: 1

sec: "C"

marks: 90

Big Data Analytics



EXERCISE:

1. Write the insert method to store the following document in MongoDB and verify. May practice crud operations.

Name : "Stephen More"

Address : {

"City" : "Banglore",

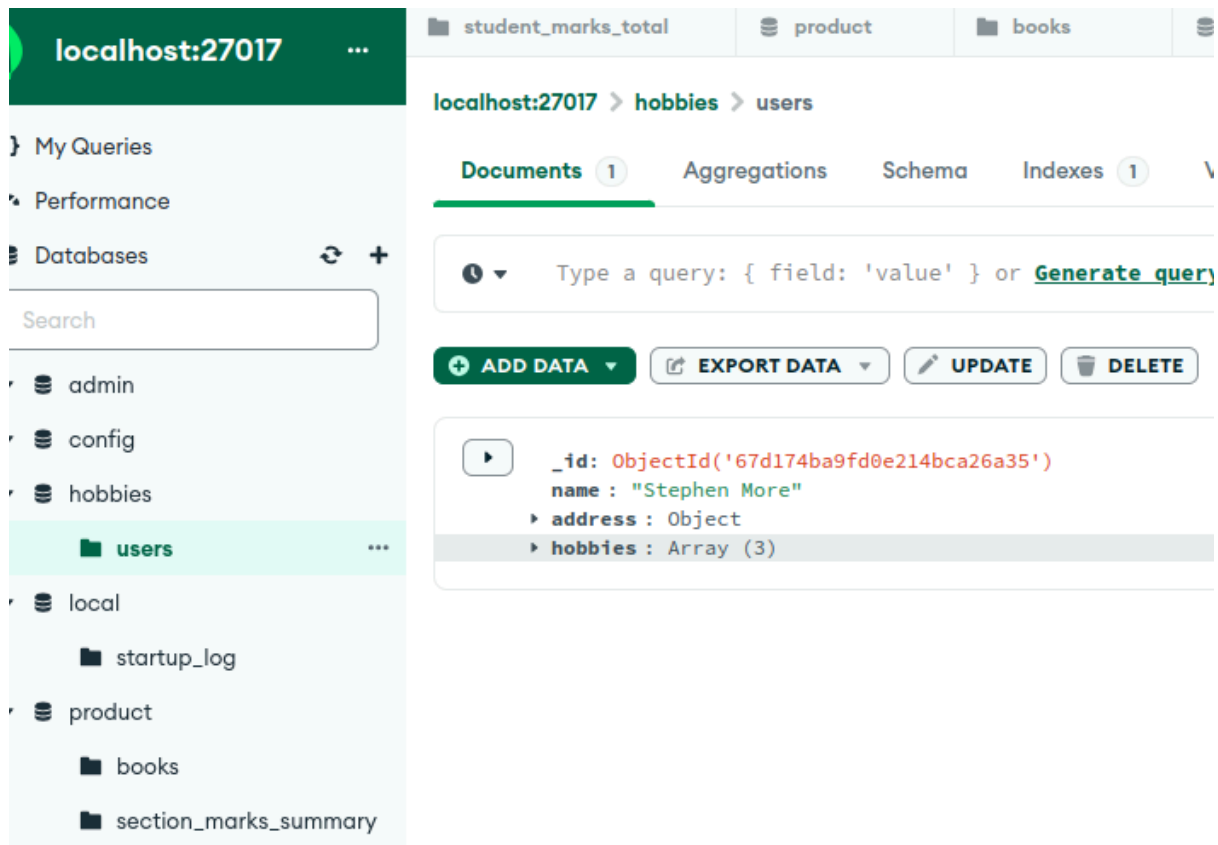
"Street" : "Electronics City",

"Affiliation" : "XYZ Ltd"

}

Hobbies : Chess, Lawn Tennis, Base Ball

```
product> use hobbies
switched to db hobbies
hobbies> db.users.insertOne({
...   name: "Stephen More",
...   address: {
...     city: "Bangalore",
...     street: "Electronics City",
...     affiliation: "XYZ Ltd"
...   },
...   hobbies: ["Chess", "Lawn Tennis", "Base Ball"]
... });
{
  acknowledged: true,
  insertedId: ObjectId('67d174ba9fd0e214bca26a35')
}
hobbies> db.users.find({ name: "Stephen More" }).pretty();
[
  {
    _id: ObjectId('67d174ba9fd0e214bca26a35'),
    name: 'Stephen More',
    address: {
      city: 'Bangalore',
      street: 'Electronics City',
      affiliation: 'XYZ Ltd'
    },
    hobbies: [ 'Chess', 'Lawn Tennis', 'Base Ball' ]
  }
]
```



2. Create a collection 'cities' as follows

```
cities> db.cities.find({}, { name: 1, population: 1, _id: 0 }).sort({ population: -1 }).pretty();
[
  { name: 'Tokyo', population: 37.4 },
  { name: 'Delhi', population: 28.514 },
  { name: 'Seoul', population: 25.674 },
  { name: 'Shanghai', population: 25.582 },
  { name: 'São Paulo', population: 21.65 },
  { name: 'Mexico City', population: 21.581 },
  { name: 'Cairo', population: 20.076 },
  { name: 'Mumbai', population: 19.98 },
  { name: 'Beijing', population: 19.618 },
  { name: 'Dhaka', population: 19.578 },
  { name: 'Osaka', population: 19.281 },
  { name: 'New York', population: 18.819 },
  { name: 'Karachi', population: 15.4 },
  { name: 'Buenos Aires', population: 14.967 },
  { name: 'Chongqing', population: 14.838 },
  { name: 'Istanbul', population: 14.751 },
  { name: 'Kolkata', population: 14.681 },
  { name: 'Manila', population: 13.482 },
  { name: 'Lagos', population: 13.463 },
  { name: 'Rio de Janeiro', population: 13.293 }
]
```

- Write a MongoDB query to find all the cities of North America.
- Write a MongoDB query to retrieve the cities names in descending order of their population.
- Write a MongoDB query to display all the cities grouped by their continent.

- Write a mongodb query to add the field “highest population”.

```
cities> db.cities.aggregate([
...   {
...     $group: {
...       _id: "$continent", // Group by continent
...       cities: { $push: "$name" } // Push city names into an array
...     }
...   }
... ]);
[
  { _id: 'Africa', cities: [ 'Lagos', 'Cairo' ] },
  {
    _id: 'South America',
    cities: [ 'Rio de Janeiro', 'São Paulo', 'Buenos Aires' ]
  },
  { _id: 'North America', cities: [ 'Mexico City', 'New York' ] },
  {
    _id: 'Asia',
    cities: [
      'Seoul', 'Mumbai',
      'Beijing', 'Shanghai',
      'Osaka', 'Tokyo',
      'Karachi', 'Dhaka',
      'Delhi', 'Kolkata',
      'Manila', 'Chongqing'
    ]
  },
  { _id: 'Europe', cities: [ 'Istanbul' ] }
]
cities> db.cities.aggregate([
...   {
...     $group: {
...       _id: null, // Group everything into a single document
...       highest_population: { $max: "$population" } // Find the highest population
...     }
...   },
...   {
...     $project: { _id: 0, highest_population: 1 } // Display only the highest_population field
...   }
... ]);
[ { highest_population: 37.4 } ]
```


The screenshot displays the MongoDB Compass interface. On the left, a sidebar shows a list of databases: admin, cities, config, hobbies, local, product, and startup_log. The 'cities' database is selected and highlighted. Below the database list, a 'Search' bar is visible. On the right, the main panel shows the 'cities' collection with 5 documents. Each document is displayed as a JSON object with fields: _id, name, country, continent, and population. The documents represent cities from South Korea, India, Nigeria, China, and Japan.

_id	name	country	continent	population
ObjectId('67d1759b9fd0e214bca26a36')	Seoul	South Korea	Asia	25.674
ObjectId('67d1759b9fd0e214bca26a37')	Mumbai	India	Asia	19.98
ObjectId('67d1759b9fd0e214bca26a38')	Lagos	Nigeria	Africa	13.463
ObjectId('67d1759b9fd0e214bca26a39')	Beijing	China	Asia	19.618
ObjectId('67d1759b9fd0e214bca26a3a')	Shanghai	China	Asia	25.582
ObjectId('67d1759b9fd0e214bca26a3b')	Osaka	Japan	Asia	19.281

3. To practice MapReduce programming in MongoDB.

Step. 3.1: Insert 5 documents as shown below in the collection named 'books'.

```
books> db.books.insertMany([
...   {
...     _id: 1,
...     Category: "Machine Learning",
...     BookName: "Machine Learning for Hackers",
...     Author: "Drew Conway",
...     qty: 25,
...     price: 400,
...     rol: 30,
...     pages: 350
...   },
...   {
...     _id: 2,
...     Category: "Business Intelligence",
...     BookName: "Fundamentals of Business Analytics",
...     Author: "Seema Acharya",
...     qty: 55,
...     price: 500,
...     rol: 30,
...     pages: 250
...   },
...   {
...     _id: 3,
...     Category: "Analytics",
...     BookName: "Competing on Analytics",
...     Author: "Thomas Davenport",
...     qty: 8,
...     price: 150,
...     rol: 20,
...     pages: 150
...   },
...   {
...     _id: 4,
...     Category: "Visualization",
...     BookName: "Visualizing Data",
...     Author: "Ben Fry",
...     qty: 12,
...     price: 325,
...     rol: 6,
...     pages: 450
...   },
...   {
...     _id: 5,
...     Category: "Web Mining",
...     BookName: "Learning R",
...     Author: "Richard Cotton",
...     qty: 5,
...     price: 850,
...     rol: 10,
...     pages: 120
...   }
... ])
```

Step. 3.2: Confirm the presence of the above documents in the “books” collection.

```
acknowledged: true,
insertedIds: { '0': 1, '1': 2, '2': 3, '3': 4, '4': 5 }
}
books> db.books.find().pretty();
[
  {
    _id: 1,
    Category: 'Machine Learning',
    BookName: 'Machine Learning for Hackers',
    Author: 'Drew Conway',
    qty: 25,
    price: 400,
    rol: 30,
    pages: 350
  },
  {
    _id: 2,
    Category: 'Business Intelligence',
    BookName: 'Fundamentals of Business Analytics',
    Author: 'Seema Acharya',
    qty: 55,
    price: 500,
    rol: 30,
    pages: 250
  },
  {
    _id: 3,
    Category: 'Analytics',
    BookName: 'Competing on Analytics',
    Author: 'Thomas Davenport',
    qty: 8,
    price: 150,
    rol: 20,
    pages: 150
  },
  {
    _id: 4,
    Category: 'Visualization',
    BookName: 'Visualizing Data',
    Author: 'Ben Fry',
    qty: 12,
    price: 325,
    rol: 6,
    pages: 450
  },
  {
    _id: 5,
    Category: 'Web Mining',
    BookName: 'Learning R',
    Author: 'Richard Cotton',
    qty: 5,
    price: 850,
    rol: 10
  }
]
```

Step. 3.3: Write map and reduce functions to split the books into the following two categories:

(a) Big Books

(b) Small Books

Books which have more than 300 pages should be in the big book category. Books which have less than 300 pages should be in the small book category.

Step. 3.4: Count the number of books in each category.

Step 3.5: Store the output as follows as documents in a new collection, called "Book_Result".

```
_id: 5,
Category: 'Web Mining',
BookName: 'Learning R',
Author: 'Richard Cotton',
qty: 5,
price: 850,
rol: 10,
pages: 120
}
]
books> var mapFunction = function () {
...   var category = this.pages > 300 ? "Big Books" : "Small Books";
...   emit(category, 1);
... };

books> var reduceFunction = function (key, values) {
...   return Array.sum(values);
... };

books> db.books.mapReduce(
...   mapFunction,
...   reduceFunction,
...   {
...     out: "Book_Result"
...   }
... );
DeprecationWarning: Collection.mapReduce() is deprecated. Use an aggregation instead.
See https://docs.mongodb.com/manual/core/map-reduce for details.
{ result: 'Book_Result', ok: 1 }
books> db.Book_Result.find().pretty();
[ { _id: 'Big Books', value: 2 }, { _id: 'Small Books', value: 3 } ]
```

Output:

The screenshot displays the MongoDB Compass web interface. On the left sidebar, the 'Databases' section is expanded, showing a list of databases: 'Student', 'admin', 'books', 'Book_Result' (highlighted in green), 'books', and 'cities'. The main panel shows the 'Book_Result' collection under the 'books' database. The 'Documents' tab is active, displaying two documents in a list view. The first document has an '_id' of 'Big Books' and a 'value' of 2. The second document has an '_id' of 'Small Books' and a 'value' of 3. The interface includes a search bar, a query input field, and buttons for 'ADD DATA', 'EXPORT DATA', and 'UPDATE'.

localhost:27017

My Queries

Performance

Databases

Search

- Student
- admin
- books
- Book_Result**
- books
- cities

localhost:27017 > books > Book_Result

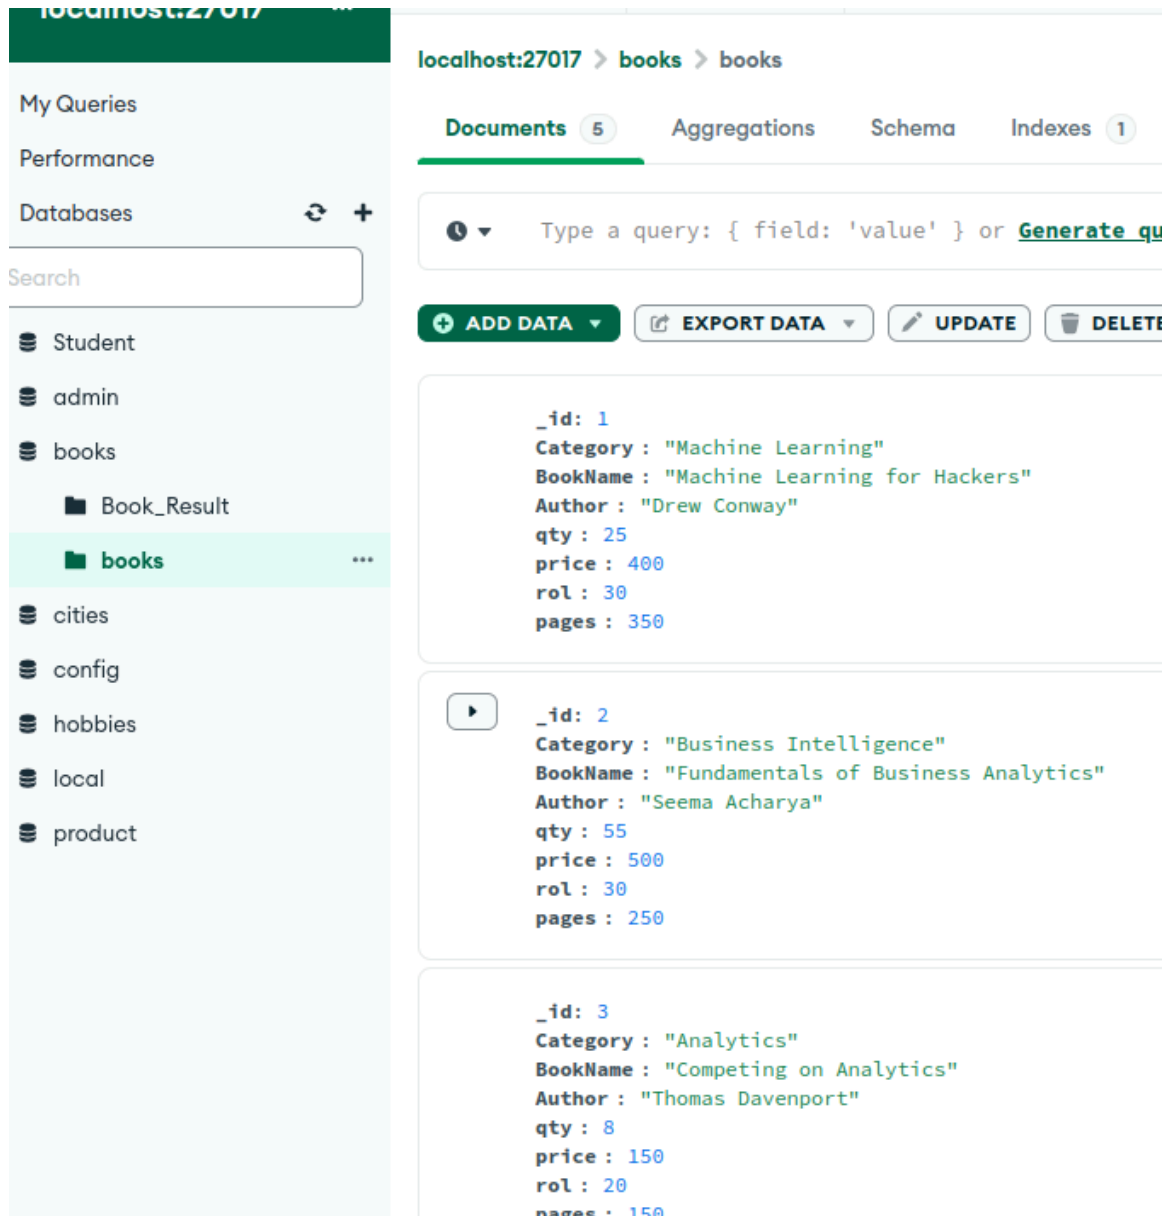
Documents 2 Aggregations Schema Indexes

Type a query: { field: 'value' } or [Generate](#)

+ ADD DATA EXPORT DATA UPDATE

```
{ "_id": "Big Books", "value": 2 }
```

```
{ "_id": "Small Books", "value": 3 }
```



The screenshot shows the MongoDB Compass web interface. On the left, a sidebar lists databases and collections. The 'local' database is selected, and the 'books' collection is highlighted. The main area shows the 'books' collection with 5 documents. Three documents are visible:

- Document 1:**

```

_id: 1
Category: "Machine Learning"
BookName: "Machine Learning for Hackers"
Author: "Drew Conway"
qty: 25
price: 400
rol: 30
pages: 350

```
- Document 2:**

```

_id: 2
Category: "Business Intelligence"
BookName: "Fundamentals of Business Analytics"
Author: "Seema Acharya"
qty: 55
price: 500
rol: 30
pages: 250

```
- Document 3:**

```

_id: 3
Category: "Analytics"
BookName: "Competing on Analytics"
Author: "Thomas Davenport"
qty: 8
price: 150
rol: 20
pages: 150

```

3. Write a mongodb query using mapreduce to find mutual friends.

-you may consider the required collections.

```

mydatabase> db.users.insertMany([
...   { "_id": 1, "name": "Alice", "friends": [2, 3, 4] },
...   { "_id": 2, "name": "Bob", "friends": [1, 3, 5] },
...   { "_id": 3, "name": "Charlie", "friends": [1, 2, 4, 5] },
...   { "_id": 4, "name": "David", "friends": [1, 3] },
...   { "_id": 5, "name": "Eve", "friends": [2, 3] }
... ])
{
  acknowledged: true,
  insertedIds: { '0': 1, '1': 2, '2': 3, '3': 4, '4': 5 }
}

```

```

mydatabase> var mapFunction = function () {
...     for (var i = 0; i < this.friends.length; i++) {
...         for (var j = i + 1; j < this.friends.length; j++) {
...             var pair = [this.friends[i], this.friends[j]].sort();
...             emit(pair.join(","), 1);
...         }
...     }
... };

mydatabase> var reduceFunction = function (key, values) {
...     return Array.sum(values);
... };

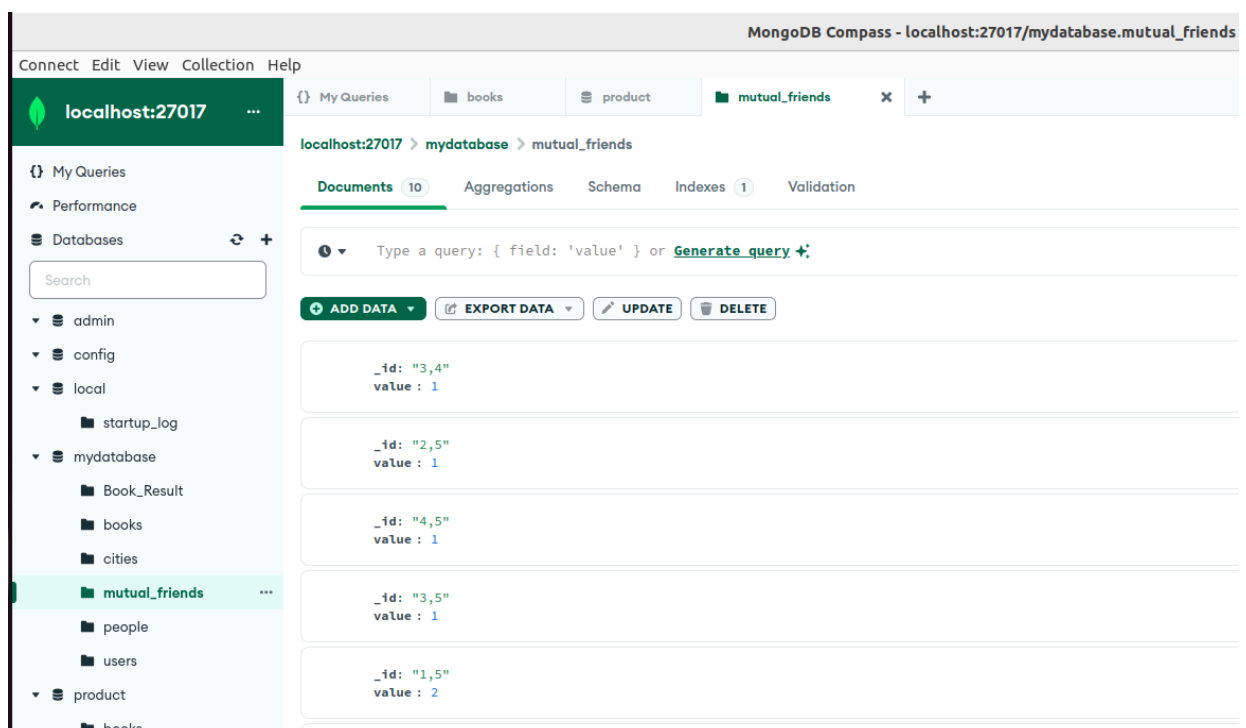
mydatabase> db.users.mapReduce(
...     mapFunction,
...     reduceFunction,
...     { out: "mutual_friends" }
... )
{ result: 'mutual_friends', ok: 1 }

```

```

mydatabase> db.mutual_friends.find().pretty()
[
  { _id: '3,4', value: 1 },
  { _id: '2,5', value: 1 },
  { _id: '4,5', value: 1 },
  { _id: '3,5', value: 1 },
  { _id: '1,5', value: 2 },
  { _id: '2,4', value: 2 },
  { _id: '2,3', value: 2 },
  { _id: '1,4', value: 1 },
  { _id: '1,3', value: 2 },
  { _id: '1,2', value: 1 }
]

```



Summarised learning:

This lab, we focused on learning MongoDB, a document-oriented NoSQL database, through hands-on exercises involving CRUD operations, queries, and advanced features like MapReduce. Key topics include connecting to MongoDB, creating and manipulating databases and collections, inserting, updating, and deleting documents, and performing queries using operators for filtering and sorting data. Additionally, the lab introduced us to MapReduce functionality for data aggregation, such as categorizing books by size and counting them. Also practical exercises for database management, data analysis, and applying MongoDB's powerful querying capabilities for real-world scenarios like mutual friend finding were implemented.