

## LAB - 06

### Neo4j

**Objective:** Students will learn how to structure data in a graph format, use Cypher queries for data manipulation, and explore real-world use cases of graph databases. The objective of this assignment is to give students practical exposure to graph databases by building a social network model using Neo4j.

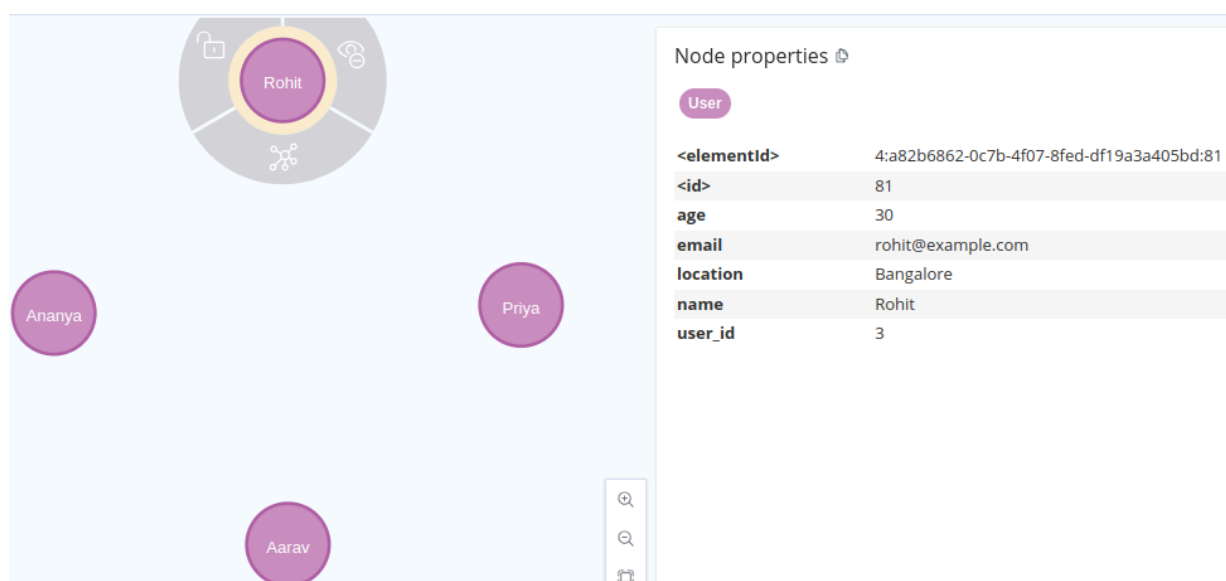
#### ASSIGNMENT:

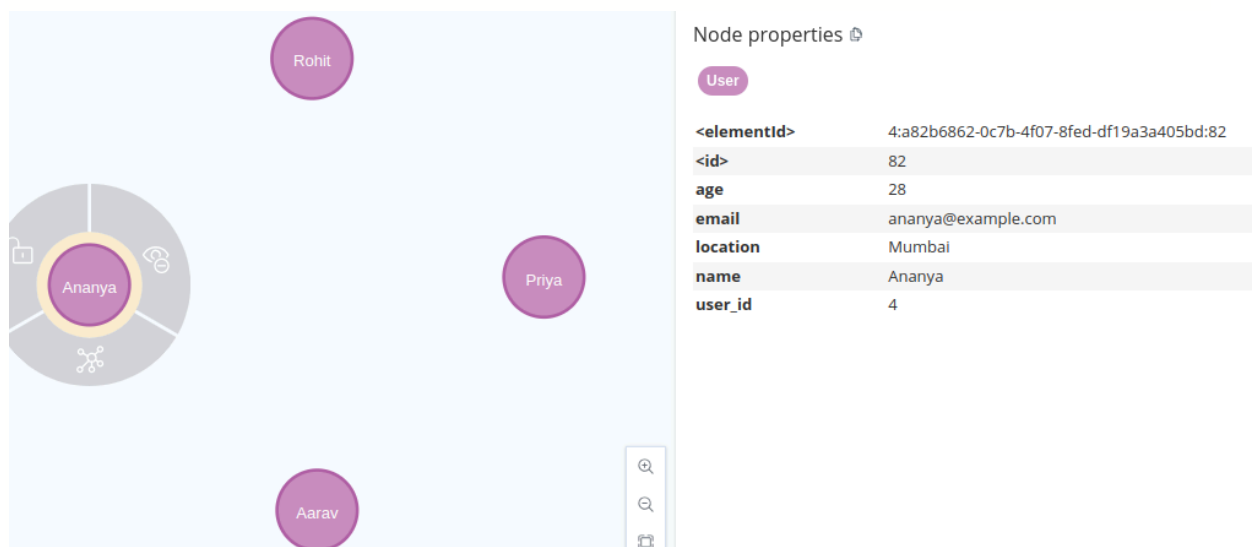
**Design the database for Social networks.**

##### 1. Nodes:

- **User Node: Represents the users of the social network.**
- **Properties: user\_id, name, age, location, email**

```
CREATE (u1:User {user_id: 1, name: "Aarav", age: 25, location: "Mumbai", email: "aarav@example.com"})
CREATE (u2:User {user_id: 2, name: "Priya", age: 27, location: "Delhi", email: "priya@example.com"})
CREATE (u3:User {user_id: 3, name: "Rohit", age: 30, location: "Bangalore", email: "rohit@example.com"})
CREATE (u4:User {user_id: 4, name: "Ananya", age: 28, location: "Mumbai", email: "ananya@example.com"})
RETURN u1, u2, u3, u4;
```





- **Post Node:** Represents a post created by a user.

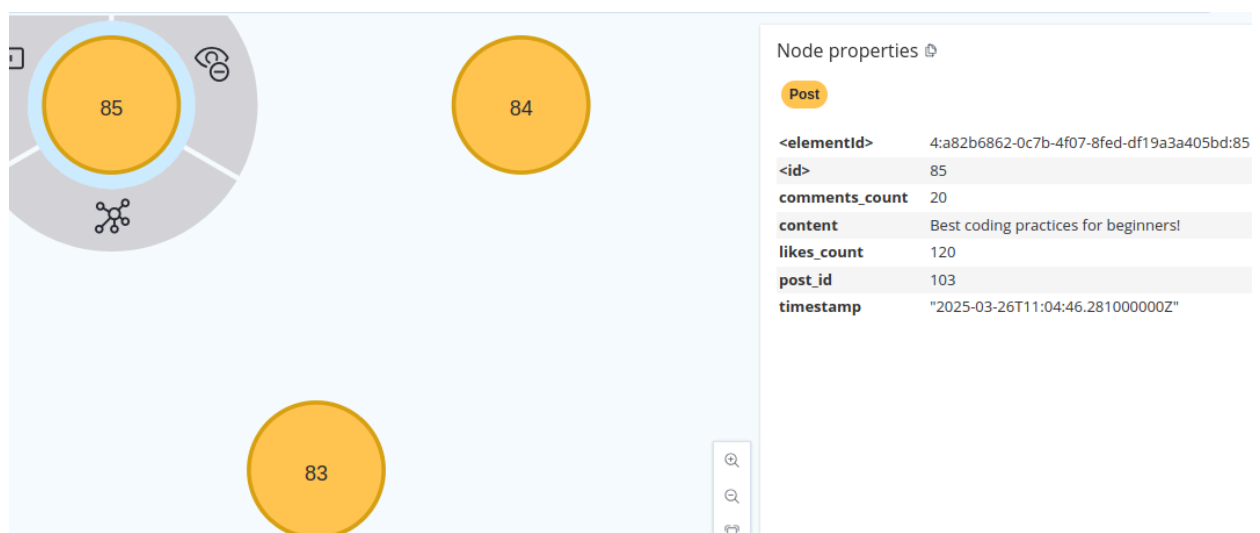
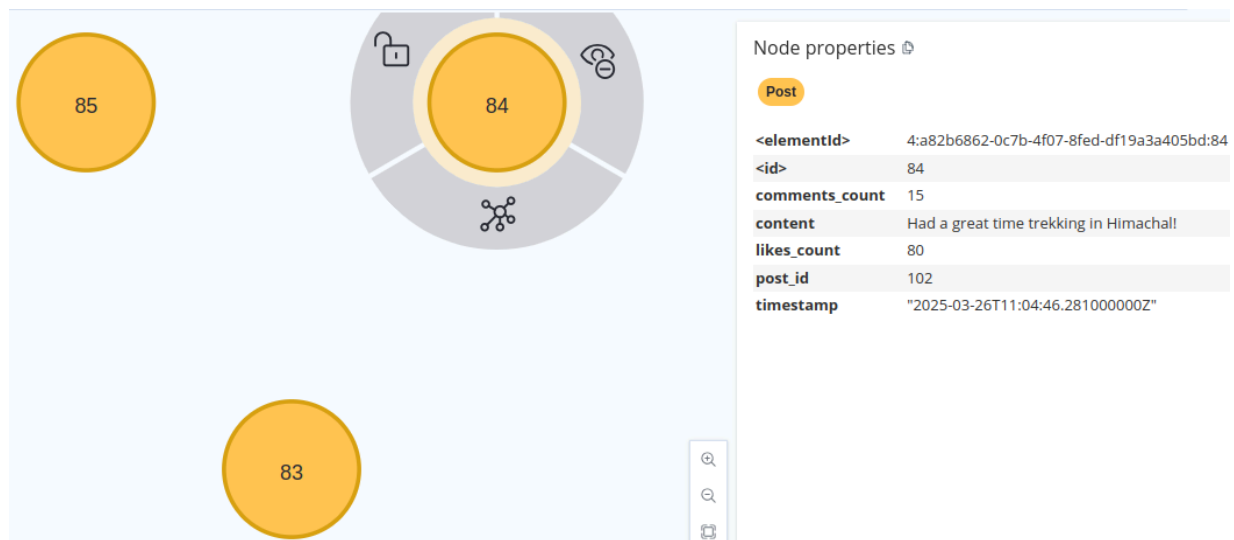
- **Properties:** `post_id`, `content`, `timestamp`, `likes_count`, `comments_count`

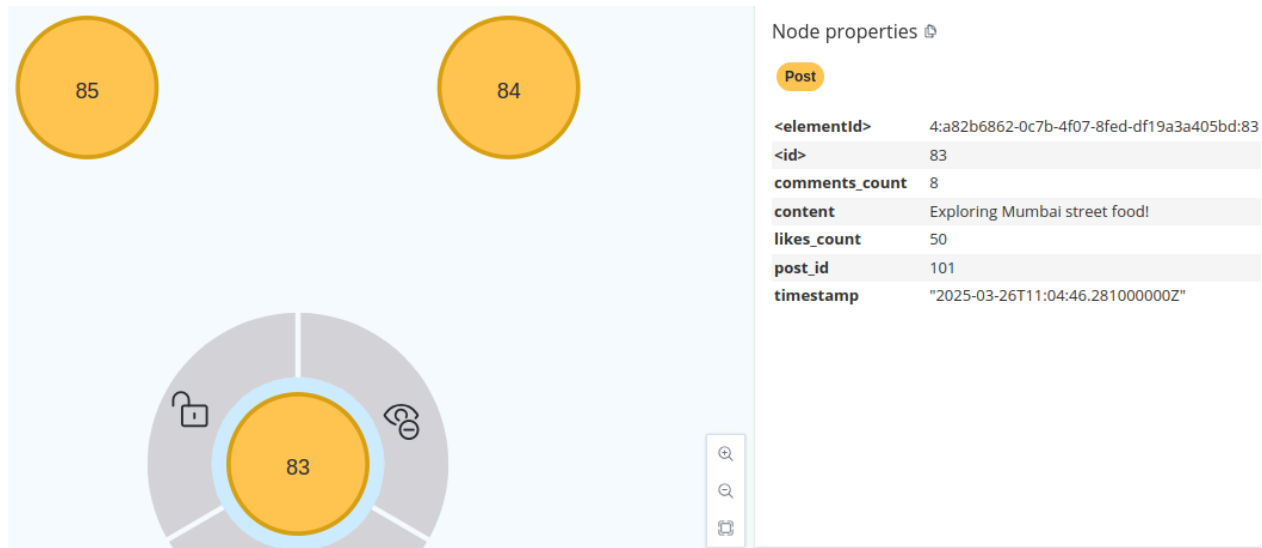
```
CREATE (p1:Post {post_id: 101, content: "Exploring Mumbai street food!", timestamp: datetime(), likes_count: 50, comments_count: 8})
```

```
CREATE (p2:Post {post_id: 102, content: "Had a great time trekking in Himachal!", timestamp: datetime(), likes_count: 80, comments_count: 15})
```

```
CREATE (p3:Post {post_id: 103, content: "Best coding practices for beginners!", timestamp: datetime(), likes_count: 120, comments_count: 20})
```

```
RETURN p1, p2, p3;
```





- **Comment Node: Represents a comment on a post.**

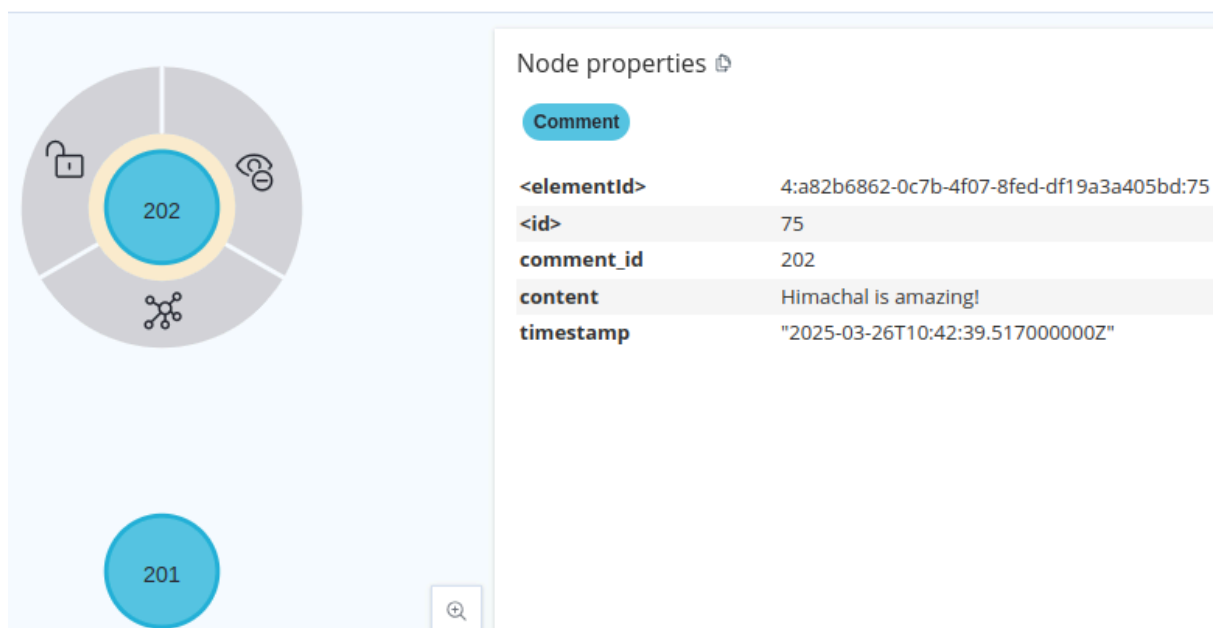
- **Properties: comment\_id, content, timestamp**

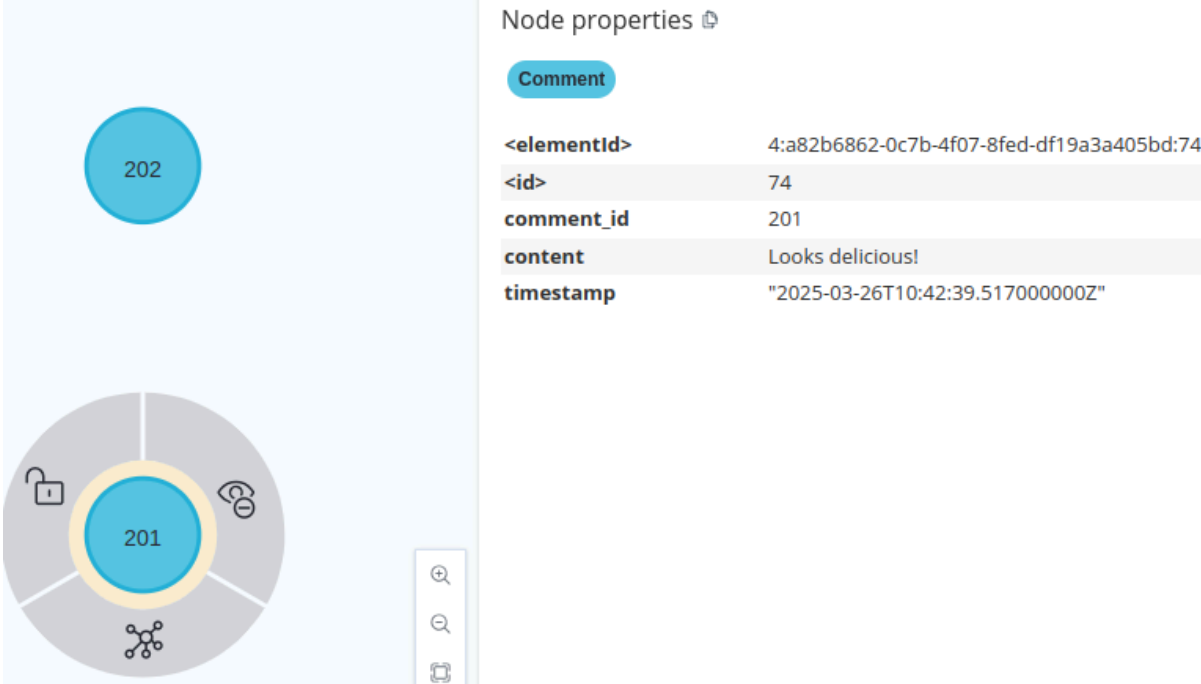
```
CREATE (c1:Comment {comment_id: 201, content: "Looks delicious!", timestamp:
datetime()})
```

```
CREATE (c2:Comment {comment_id: 202, content: "Himachal is amazing!", timestamp:
datetime()})
```

```
RETURN c1, c2;
```

```
ip: datetime()})
```





Node properties

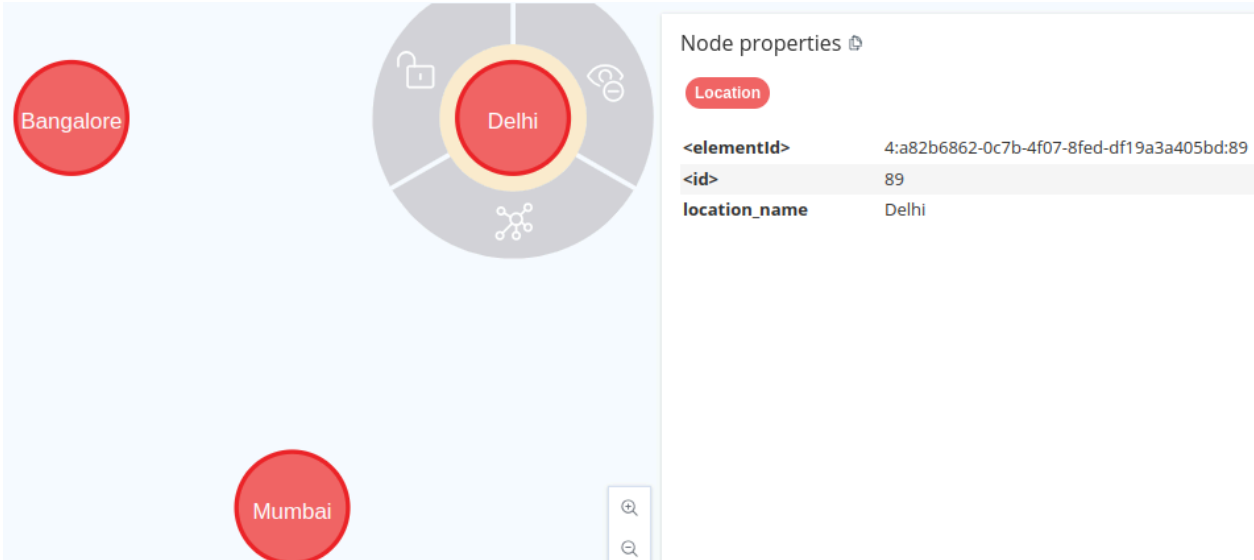
Comment

|             |   |
|-------------|---|
| <elementId> | 4:a82b6862-0c7b-4f07-8fed-df19a3a405bd:74 |
| <id>        | 74  |
| comment_id  | 201                                       |
| content     | Looks delicious!                          |
| timestamp   | "2025-03-26T10:42:39.517000000Z"          |

● **Location Node: (Optional)** If you're going to query based on location.

- **Properties:** location\_name

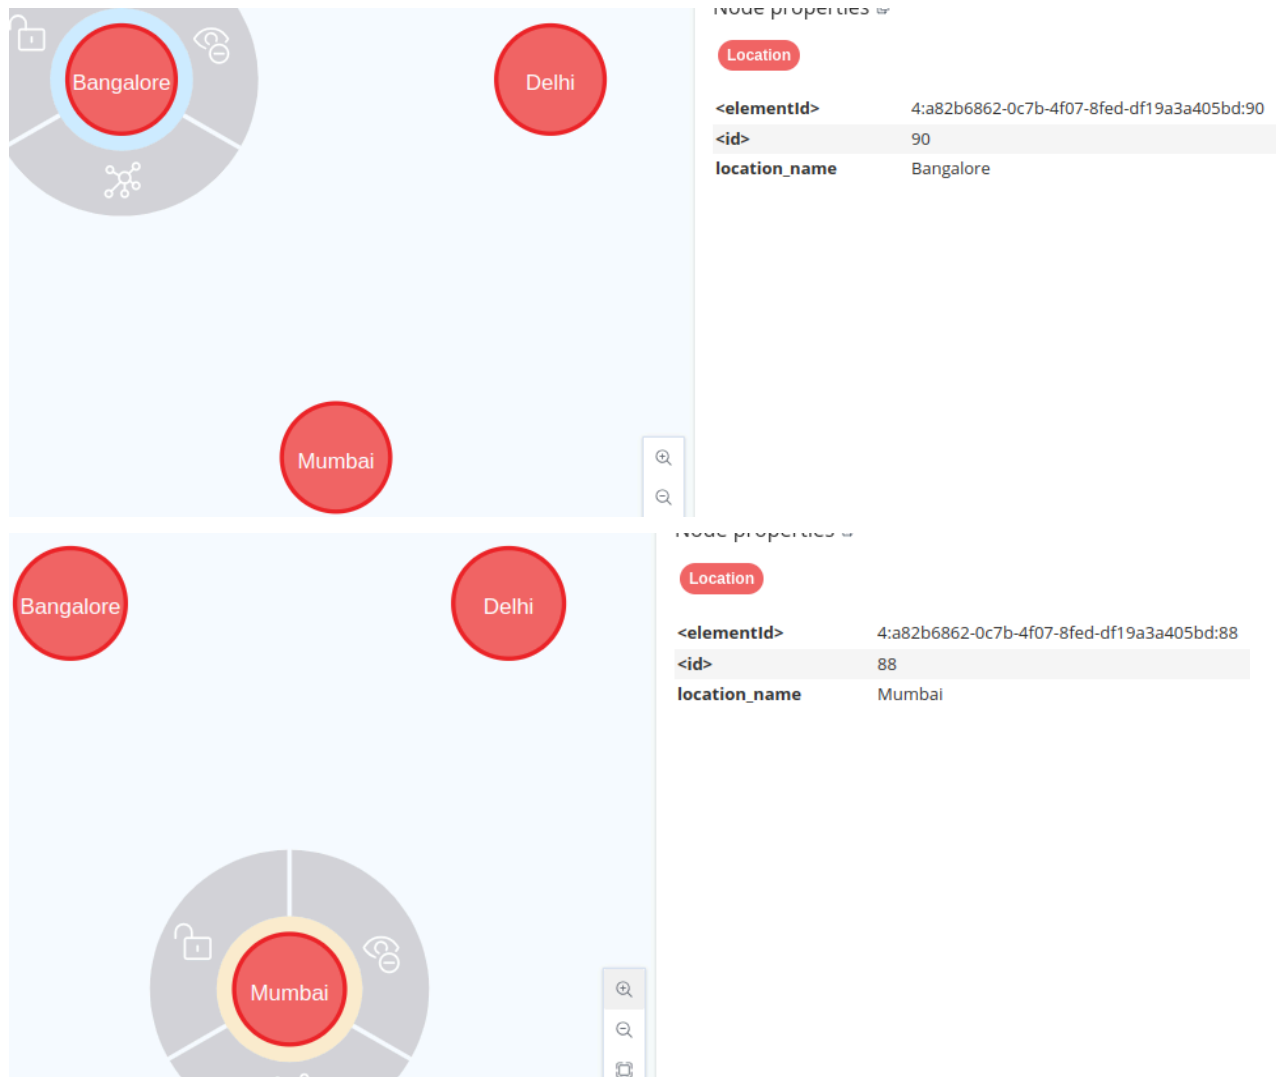
```
CREATE (l1:Location {location_name: "Mumbai"})
CREATE (l2:Location {location_name: "Delhi"})
CREATE (l3:Location {location_name: "Bangalore"})
RETURN l1, l2, l3;
```



Node properties

Location

|               |   |
|---------------|---|
| <elementId>   | 4:a82b6862-0c7b-4f07-8fed-df19a3a405bd:89 |
| <id>          | 89  |
| location_name | Delhi                                     |



## 2. Relationships:

- **FRIENDS\_WITH:** Represents the friendship relationship between two users.
- **Properties:** since (the date when the friendship started)

```
MATCH (a:User {name: "Aarav"}), (b:User {name: "Priya"})
CREATE (a)-[:FRIENDS_WITH {since: date("2022-06-15")}]>(b);
```

```
MATCH (b:User {name: "Priya"}), (c:User {name: "Rohit"})
CREATE (b)-[:FRIENDS_WITH {since: date("2021-03-10")}]>(c);
```

```
MATCH (a:User {name: "Aarav"}), (d:User {name: "Ananya"})
CREATE (a)-[:FRIENDS_WITH {since: date("2023-01-25")}]>(d);
```

```
neo4j$ MATCH (a:User {name: "Aarav"}), (b:User {name: "Priya"}) CREATE (a)-[:FRIENDS_WITH {since: date("2022-06-15")}]>(b)
neo4j$ MATCH (b:User {name: "Priya"}), (c:User {name: "Rohit"}) CREATE (b)-[:FRIENDS_WITH {since: date("2021-03-10")}]>(c)
neo4j$ MATCH (a:User {name: "Aarav"}), (d:User {name: "Ananya"}) CREATE (a)-[:FRIENDS_WITH {since: date("2023-01-25")}]>(d)
```

- **CREATED:** Represents a user creating a post.

- **Properties:** created\_at (timestamp when the post was created)

```
MATCH (u:User {name: "Aarav"}), (p:Post {post_id: 101})
```

```
CREATE (u)-[:CREATED {created_at: datetime()}]->(p);
```

```
MATCH (u:User {name: "Priya"}), (p:Post {post_id: 102})
```

```
CREATE (u)-[:CREATED {created_at: datetime()}]->(p);
```

```
neo4j$ MATCH (u:User {name: "Aarav"}), (p:Post {post_id: 101}) CREATE (u)-[:CREATED {created_at: datetime()}]->(p)
neo4j$ MATCH (u:User {name: "Priya"}), (p:Post {post_id: 102}) CREATE (u)-[:CREATED {created_at: datetime()}]->(p)
```

- **COMMENTED:** Represents a user commenting on a post.

- **Properties:** commented\_at (timestamp when the comment was made)

```
MATCH (u:User {name: "Rohit"}), (c:Comment {comment_id: 201}), (p:Post {post_id: 101})
```

```
CREATE (u)-[:COMMENTED {commented_at: datetime()}]->(c)-[:ON_POST]->(p);
```

```
MATCH (u:User {name: "Ananya"}), (c:Comment {comment_id: 202}), (p:Post {post_id: 102})
```

```
CREATE (u)-[:COMMENTED {commented_at: datetime()}]->(c)-[:ON_POST]->(p);
```

```
neo4j$ MATCH (u:User {name: "Rohit"}), (c:Comment {comment_id: 201}), (p:Post {post_id: 101}) CREATE (u)-[:COMMENTED {commented_at: datetime()}]->(c)-[:ON_POST]->(p)
neo4j$ MATCH (u:User {name: "Ananya"}), (c:Comment {comment_id: 202}), (p:Post {post_id: 102}) CREATE (u)-[:COMMENTED {commented_at: datetime()}]->(c)-[:ON_POST]->(p)
```

**SUCCESS**

Set 126 properties, created 252 relationships, completed after 6 ms.

- **LOCATED\_IN:** (Optional) Connects a user to a location.

- **Properties:** since (timestamp when the user joined this location)

```
MATCH (u:User {name: "Aarav"}), (l:Location {location_name: "Mumbai"})
```

```
CREATE (u)-[:LOCATED_IN {since: date("2020-08-10")}]>(l);
```

```
MATCH (u:User {name: "Rohit"}), (l:Location {location_name: "Bangalore"})
```

```
CREATE (u)-[:LOCATED_IN {since: date("2021-05-20")}]>(l);
```

```
neo4j$ MATCH (u:User {name: "Aarav"}), (l:Location {location_name: "Mumbai"}) CREATE (u)-[:LOCATED_IN {since: date("2020-08-10")}]->(l)
neo4j$ MATCH (u:User {name: "Rohit"}), (l:Location {location_name: "Bangalore"}) CREATE (u)-[:LOCATED_IN {since: date("2021-05-20")}]->(l)
```

### 3. Queries:

#### a. Retrieve all friends of a specific user.

```
MATCH (:User {name: "Aarav"})-[:FRIENDS_WITH]-(friend)
RETURN friend.name;
```

|   | friend.name |
|---|-------------|
| 1 | "Priya"     |
| 2 | "Priya"     |
| 3 | "Priya"     |
| 4 | "Ananya"    |
| 5 | "Ananya"    |
| 6 | "Ananya"    |

#### b. Find the most popular post (by the number of comments or likes).

```
MATCH (p:Post)
RETURN p.content, p.likes_count, p.comments_count
ORDER BY p.likes_count DESC, p.comments_count DESC
LIMIT 1;
```

|  | p.content                              | p.likes_count | p.comments_count |
|--|--|---------------|------------------|
|  | "Best coding practices for beginners!" | 120           | 20               |



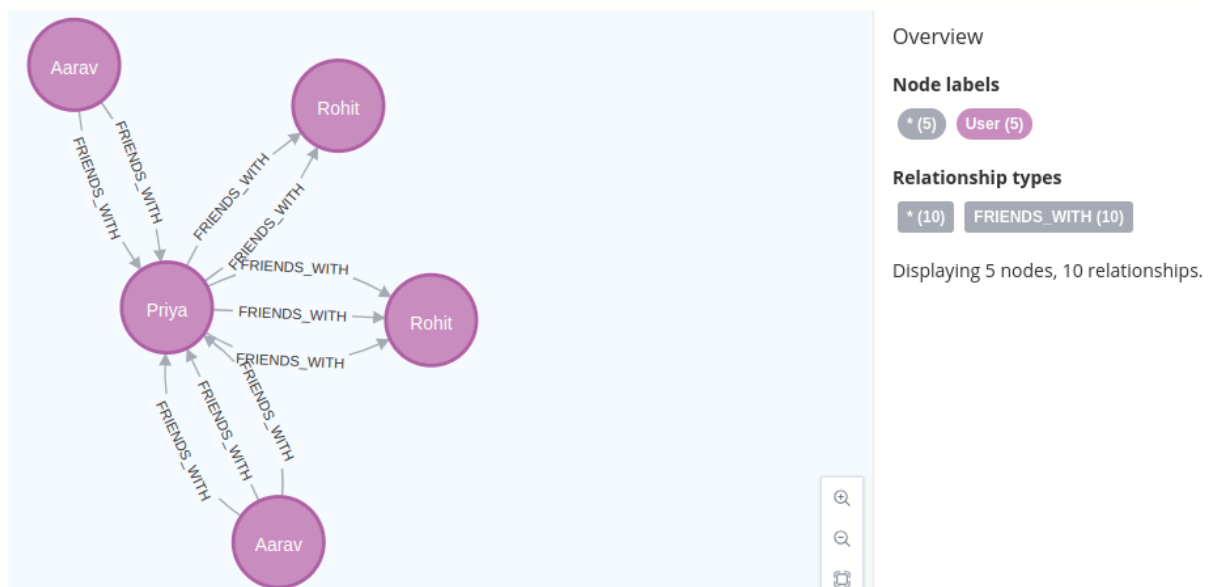
**c. List all posts by a specific user and their associated comments.**

```
MATCH (u:User {name: "Priya"})-[:CREATED]->(p:Post)
OPTIONAL MATCH (p)-[:ON_POST]-(c:Comment)
RETURN p.content, collect(DISTINCT c.content) AS comments;
```

| p.content                                | comments                                     |
|--|--|
| "Neo4j is awesome!"                      | ["I totally agree!", "Himachal is amazing!"] |
| "Had a great time trekking in Himachal!" | ["I totally agree!", "Himachal is amazing!"] |

**d. Identify the shortest path between two users (mutual friends or common connections).**

```
MATCH path = shortestPath((u1:User {name: "Aarav"})-[:FRIENDS_WITH*]-(u2:User {name: "Rohit"}))
RETURN path;
```



**e. List all posts made by users who are located in a specific region.**

```
MATCH (u:User)-[:LOCATED_IN]->(l:Location {location_name: "Mumbai"}),
(u)-[:CREATED]->(p:Post)
RETURN DISTINCT u.name, p.content;
```

|   | u.name  | p.content                       |
|---|---------|---------------------------------|
| 1 | "Aarav" | "Hello World!"                  |
| 2 | "Aarav" | "Exploring Mumbai street food!" |

**f. Find users who have commented on more than 5 posts in the last month.**

```

MATCH (u:User)-[:COMMENTED]->(c:Comment)
WHERE c.timestamp >= datetime() - duration({months: 1})
WITH u, count(DISTINCT c) AS commentCount
WHERE commentCount > 5
RETURN u.name, commentCount;

```

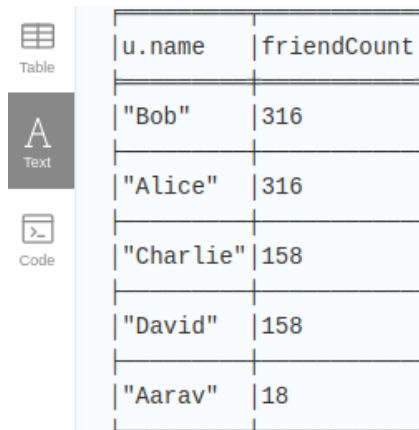
| u.name   | commentCount |
|----------|--------------|
| "Rohit"  | 7            |
| "Rohit"  | 7            |
| "Ananya" | 7            |
| "Ananya" | 7            |

**g. Find users who are the most connected (central users in the graph).**

```

MATCH (u:User)-[:FRIENDS_WITH]-(friend)
RETURN u.name, count(friend) AS friendCount
ORDER BY friendCount DESC
LIMIT 5;

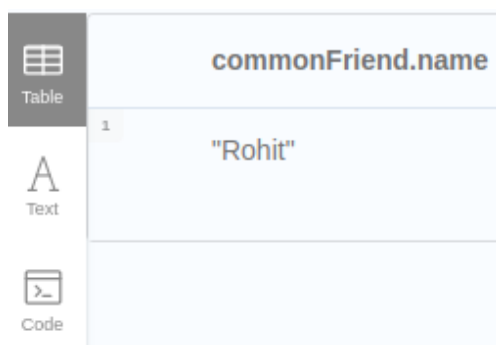
```



| u.name    | friendCount |
|-----------|-------------|
| "Bob"     | 316         |
| "Alice"   | 316         |
| "Charlie" | 158         |
| "David"   | 158         |
| "Aarav"   | 18          |

#### h. Find mutual friends between two users.

```
MATCH (u1:User {name:
"Aarav"})-[:FRIENDS_WITH]-(commonFriend)-[:FRIENDS_WITH]-(u2:User {name:
"Priya"})
RETURN commonFriend.name;
```



| commonFriend.name |
|-------------------|
| "Rohit"           |

#### Summarised learning:

In this lab, I worked with Neo4j to build a social network model, incorporating nodes such as users, posts, comments, and locations, and relationships like friendship, post creation, and commenting. I wrote and executed various Cypher queries to explore the relationships and interactions within the network. The tasks included retrieving friends of a specific user, identifying the most popular posts, listing posts with comments, finding mutual friends, and analyzing user activity based on location and comment counts. This lab helped me gain hands-on experience in structuring data in a graph format and efficiently querying it to extract valuable insights.

## Cassandra

**Aim:- To understand and implement the use of distributed NoSQL database/s like Cassandra.**

**1. Objective : Students will get familiar with the distributed database. they are able to understand the architecture perspective of cassandra along with Cassandra Query Language(CQL).**

Dashboard / Serverless Databases

**MyCassandraProject01**

Vector

Pending...

c60aa37d-8fc3-4b9f-9f5a-4063af6695eb

Overview

Data Explorer

Settings



**Your database is initializing..**

Initializing takes a few minutes

To find/describe the keyspaces

```
Connected as isha130612@gmail.com.
Connected to cndb at cassandra.ingress:9042.
[cqlsh 6.8.0 | Cassandra 4.0.0.6816 | CQL spec 3.4.5 | Native protocol v4 | TLS]
Use HELP for help.
token@cqlsh> DESCRIBE KEYSPACES

system_auth          system               datastax_sla
system_schema        system_traces        system_views
data_endpoint_auth   default_keyspace     system_virtual_schema
```

To use the keyspace

```
token@cqlsh> USE default_keyspace;
```

To create a table in keyspace use CREATE TABLE COMMAND

```
token@cqlsh:default_keyspace> CREATE TABLE EMPLOYEE (empId int PRIMARY KEY, name text, departme
nt text);
token@cqlsh:default_keyspace> []
```

To display the list of tables present in keyspace

```
token@cqlsh:default_keyspace> DESCRIBE TABLES;
```

```
employee
```

```
token@cqlsh:default_keyspace> describe tables;
```

```
employee student
```

To display the structure of a table use describe table tablename

```
token@cqlsh:default_keyspace> describe tables;
```

```
employee student
```

```
token@cqlsh:default_keyspace> describe table student;
```

```
CREATE TABLE default_keyspace.student (
  stud_id int PRIMARY KEY,
  result float,
  sem int,
  stud_name text
) WITH additional_write_policy = '99p'
AND bloom_filter_fp_chance = 0.01
AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
AND comment = ''
AND compaction = {'class': 'org.apache.cassandra.db.compaction.UnifiedCompactionStrategy'}
AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
AND crc_check_chance = 1.0
AND default_time_to_live = 0
AND gc_grace_seconds = 864000
AND max_index_interval = 2048
AND memtable_flush_period_in_ms = 0
AND min_index_interval = 128
AND read_repair = 'BLOCKING'
AND speculative_retry = '99p';
```

## Upsert Operation in Cassandra

Cassandra handles upserts with its INSERT and UPDATE statements, where:

- **INSERT Statement:** If a record with the specified primary key does not exist, INSERT creates a new record. If it does exist, INSERT overwrites the existing record with the new values.
- **UPDATE Statement:** If the primary key exists, UPDATE modifies the existing record with the new values. If it does not exist, UPDATE effectively performs an upsert by creating a new record with the specified values.

```
token@cqlsh:default_keyspace> INSERT INTO student(stud_id, stud_name, sem, result)
... VALUES (1, 'Shital', 5, 95.8);
token@cqlsh:default_keyspace> INSERT INTO student(stud_id, stud_name, sem, result)
... VALUES (2, 'John', 7, 60.8);
token@cqlsh:default_keyspace> INSERT INTO student(stud_id, stud_name, sem, result)
... VALUES (3, 'Jiya', 7, 80.08);
token@cqlsh:default_keyspace> INSERT INTO student(stud_id, stud_name, sem, result)
... VALUES (4, 'Ram', 5, 60);
```

To display the records use select

```
token@cqlsh:default_keyspace> SELECT * FROM student;
```

| stud_id | result | sem | stud_name |
|---------|--------|-----|-----------|
| 1       | 95.8   | 5   | Shital    |
| 2       | 60.8   | 7   | John      |
| 4       | 60     | 5   | Ram       |
| 3       | 80.08  | 7   | Jiya      |

(4 rows)

### Query Example: (Student Result)

```
token@cqlsh:default_keyspace> INSERT INTO student(stud_id, stud_name, sem, result)
... VALUES (5, 'Gita', null, 85.5);
token@cqlsh:default_keyspace> UPDATE student
... SET stud_name = 'Jerry'
... WHERE stud_id = 2;
token@cqlsh:default_keyspace> SELECT * FROM student;
```

| stud_id | result | sem  | stud_name |
|---------|--------|------|-----------|
| 5       | 85.5   | null | Gita      |
| 1       | 95.8   | 5    | Shital    |
| 2       | 60.8   | 7    | Jerry     |
| 4       | 60     | 5    | Ram       |
| 3       | 80.08  | 7    | Jiya      |

(5 rows)

Update will work as an insert if data is not present

```
token@cqlsh:default_keyspace> UPDATE student
... SET stud_name = 'Mitul'
... WHERE stud_id = 3;
token@cqlsh:default_keyspace> SELECT * FROM student;
```

| stud_id | result | sem  | stud_name |
|---------|--------|------|-----------|
| 5       | 85.5   | null | Gita      |
| 1       | 95.8   | 5    | Shital    |
| 2       | 60.8   | 7    | Jerry     |
| 4       | 60     | 5    | Ram       |
| 3       | 80.08  | 7    | Mitul     |

(5 rows)

Delete operation

```
token@cqlsh:default_keyspace> DELETE FROM student
... WHERE stud_id = 5;
token@cqlsh:default_keyspace> SELECT * FROM student;
```

| stud_id | result | sem | stud_name |
|---------|--------|-----|-----------|
| 1       | 95.8   | 5   | Shital    |
| 2       | 60.8   | 7   | Jerry     |
| 4       | 60     | 5   | Ram       |
| 3       | 80.08  | 7   | Mitul     |

(4 rows)

## Design 1: By Employee ID and Department

After creating table, we insert these data values:

```
token@cqlsh:default_keyspace> INSERT INTO employee_by_id_department (employee_id, department, name, position, salary)
... VALUES (101, 'Engineering', 'Alice Smith', 'Software Engineer', 85000.00);
token@cqlsh:default_keyspace> INSERT INTO employee_by_id_department (employee_id, department, name, position, salary)
... VALUES (102, 'Engineering', 'Bob Johnson', 'DevOps Engineer', 90000.00);
token@cqlsh:default_keyspace> INSERT INTO employee_by_id_department (employee_id, department, name, position, salary)
... VALUES (103, 'HR', 'Clara Adams', 'HR Manager', 75000.00);
token@cqlsh:default_keyspace> INSERT INTO employee_by_id_department (employee_id, department, name, position, salary)
... VALUES (104, 'Finance', 'David Brown', 'Financial Analyst', 80000.00);
token@cqlsh:default_keyspace> INSERT INTO employee_by_id_department (employee_id, department, name, position, salary)
... VALUES (105, 'Sales', 'Emma Wilson', 'Sales Manager', 95000.00);
token@cqlsh:default_keyspace> INSERT INTO employee_by_id_department (employee_id, department, name, position, salary)
... VALUES (106, 'IT', 'Frank Taylor', 'IT Administrator', 78000.00);
token@cqlsh:default_keyspace> SELECT * FROM employee_by_id_department;
```

Output:

```
token@cqlsh:default_keyspace> SELECT * FROM employee_by_id_department;
```

| employee_id | department  | salary | name         | position          |
|-------------|-------------|--------|--------------|-------------------|
| 104         | Finance     | 80000  | David Brown  | Financial Analyst |
| 105         | Sales       | 95000  | Emma Wilson  | Sales Manager     |
| 102         | Engineering | 90000  | Bob Johnson  | DevOps Engineer   |
| 103         | HR          | 75000  | Clara Adams  | HR Manager        |
| 106         | IT          | 78000  | Frank Taylor | IT Administrator  |
| 101         | Engineering | 85000  | Alice Smith  | Software Engineer |

(6 rows)

```
token@cqlsh:default_keyspace> SELECT * FROM employee_by_id_department
... WHERE employee_id = 101 AND department = 'Engineering';
```

| employee_id | department  | salary | name        | position          |
|-------------|-------------|--------|-------------|-------------------|
| 101         | Engineering | 85000  | Alice Smith | Software Engineer |

(1 rows)

## Design 2: By Department and Position

```

token@cqlsh:default_keyspace> CREATE TABLE employee_by_department_position (
... department text,
... position text,
... employee_id int,
... name text,
... salary float,
... PRIMARY KEY ((department, position), employee_id)
... );
token@cqlsh:default_keyspace>
token@cqlsh:default_keyspace> INSERT INTO employee_by_department_position (department, position, employee_id, name, salary)
... VALUES ('Engineering', 'Software Engineer', 101, 'Alice Smith', 85000.00);
token@cqlsh:default_keyspace>
token@cqlsh:default_keyspace> INSERT INTO employee_by_department_position (department, position, employee_id, name, salary)
... VALUES ('Engineering', 'Senior Developer', 102, 'Bob Johnson', 95000.00);
token@cqlsh:default_keyspace>
token@cqlsh:default_keyspace> INSERT INTO employee_by_department_position (department, position, employee_id, name, salary)
... VALUES ('Engineering', 'Engineering Manager', 103, 'Charlie Brown', 120000.00);
token@cqlsh:default_keyspace>
token@cqlsh:default_keyspace> INSERT INTO employee_by_department_position (department, position, employee_id, name, salary)
... VALUES ('Sales', 'Sales Executive', 201, 'David Wilson', 75000.00);
token@cqlsh:default_keyspace>
token@cqlsh:default_keyspace> INSERT INTO employee_by_department_position (department, position, employee_id, name, salary)
... VALUES ('Sales', 'Sales Manager', 202, 'Eve Davis', 90000.00);
token@cqlsh:default_keyspace>
token@cqlsh:default_keyspace> INSERT INTO employee_by_department_position (department, position, employee_id, name, salary)
... VALUES ('Marketing', 'Marketing Specialist', 301, 'Frank Miller', 70000.00);
token@cqlsh:default_keyspace>
token@cqlsh:default_keyspace> INSERT INTO employee_by_department_position (department, position, employee_id, name, salary)
... VALUES ('Marketing', 'Marketing Manager', 302, 'Grace Lee', 95000.00);

```

Query Example:

Query by Department and Position

```

token@cqlsh:default_keyspace> SELECT * FROM employee_by_department_position WHERE department = 'Engineering' AND position = 'Senior Developer';

```

| department  | position         | employee_id | name        | salary |
|-------------|------------------|-------------|-------------|--------|
| Engineering | Senior Developer | 102         | Bob Johnson | 95000  |

(1 rows)

```

token@cqlsh:default_keyspace> SELECT * FROM employee_by_department_position
... WHERE department = 'Sales'
... ALLOW FILTERING ;

```

| department | position        | employee_id | name         | salary |
|------------|-----------------|-------------|--------------|--------|
| Sales      | Sales Executive | 201         | David Wilson | 75000  |
| Sales      | Sales Manager   | 202         | Eve Davis    | 90000  |

(2 rows)

```

token@cqlsh:default_keyspace> SELECT * FROM employee_by_department_position
... WHERE department = 'Engineering'
... ALLOW FILTERING ;

```

| department  | position            | employee_id | name          | salary  |
|-------------|---------------------|-------------|---------------|---------|
| Engineering | Software Engineer   | 101         | Alice Smith   | 85000   |
| Engineering | Senior Developer    | 102         | Bob Johnson   | 95000   |
| Engineering | Engineering Manager | 103         | Charlie Brown | 1.2e+05 |

(3 rows)



**Query example: datatypes uuid,timestamp (Table user)**

```
token@cqlsh:default_keyspace> CREATE TABLE user(
...     id uuid,
...     name text,
...     pincode int,
...     noofpost counter,
...     dateofpost timestamp,
...     PRIMARY KEY(id)
... );
```

**AlreadyExists: Table 'default\_keyspace.user' already exists**

```
token@cqlsh:default_keyspace> INSERT INTO user(id, name, pincode, dateofpost)
...     VALUES (uuid(), 'Shital', 123, '2011-02-03 04:05+0000');
token@cqlsh:default_keyspace> INSERT INTO user(id, name, pincode, dateofpost)
...     VALUES (uuid(), 'Maria', 456, '2012-03-03 04:05+0000');
token@cqlsh:default_keyspace> INSERT INTO user(id, name, pincode, dateofpost)
...     VALUES (uuid(), 'John', 888, '2022-02-03 04:05+0000');
token@cqlsh:default_keyspace> SELECT * FROM user;
```

| id                                   | dateofpost                      | name   | pincode |
|--------------------------------------|---------------------------------|--------|---------|
| e18f424b-e073-4b64-9df8-bccff52d0d56 | 2022-02-03 04:05:00.000000+0000 | John   | 888     |
| f0623668-bc91-4bb6-9c71-de6cbd64e8b8 | 2012-03-03 04:05:00.000000+0000 | Maria  | 456     |
| 959b9365-57e7-4c97-9a56-e8cf3dbca3ec | 2011-02-03 04:05:00.000000+0000 | Shital | 123     |
| 300b4420-84ad-40f5-b2d8-e24dad419323 | 2022-02-03 04:05:00.000000+0000 | John   | 888     |
| 00460618-12b3-472c-8547-6f22d7c638c0 | 2011-02-03 04:05:00.000000+0000 | Shital | 123     |
| 753f9c32-30ce-49b8-97da-a75de721def7 | 2012-03-03 04:05:00.000000+0000 | Maria  | 456     |

(6 rows)

To display output

```
token@cqlsh:default_keyspace> SELECT * FROM student WHERE stud_id = 2;
```

| stud_id | result | sem | stud_name |
|---------|--------|-----|-----------|
| 2       | 60.8   | 7   | Jerry     |

(1 rows)

**COUNTER IN CASSANDRA:****1. Create a Counter Table**

In Cassandra, counter columns require a PRIMARY KEY, and counters cannot exist with non-counter columns except the primary key.

```
token@cqlsh:default_keyspace> CREATE TABLE page_views (
...     page_id int PRIMARY KEY,
...     view_count counter
... );
```

## 2. Increment the Counter

```
token@cqlsh:default_keyspace> UPDATE page_views
...     SET view_count = view_count + 1
...     WHERE page_id = 1;
token@cqlsh:default_keyspace> UPDATE page_views
...     SET view_count = view_count + 5
...     WHERE page_id = 1;
```

## 3. Select Data from the Counter Table

```
token@cqlsh:default_keyspace> SELECT * FROM page_views;

page_id | view_count
-----+-----
      1 |         6
(1 rows)
```

## 4. Decrement the Counter

```
token@cqlsh:default_keyspace> UPDATE page_views
...     SET view_count = view_count - 2
...     WHERE page_id = 1;
token@cqlsh:default_keyspace> SELECT * FROM page_views;

page_id | view_count
-----+-----
      1 |         4
(1 rows)
```

## 5. Insert More Pages and Track Views

```
token@cqlsh:default_keyspace> UPDATE page_views
...     SET view_count = view_count + 3
...     WHERE page_id = 2;
token@cqlsh:default_keyspace> UPDATE page_views
...     SET view_count = view_count + 10
...     WHERE page_id = 3;
token@cqlsh:default_keyspace> SELECT * FROM page_views;

page_id | view_count
-----+-----
      1 |         4
      2 |         3
      3 |        10
(3 rows)
```

This completes all the steps for **creating, incrementing, decrementing, and deleting** counter values in Cassandra

**EXERCISE:****1. Create table account as follows**

```
token@cqlsh:default_keyspace> create table account(
    ... accno int,
    ... custid int,
    ... balance int,
    ... aod date,
    ... atype text,
    ... astatus text,
    ... PRIMARY KEY(atype, accno));
```

**2. Insert records into the table account as follows.(using upsert operations in cassandra)**

```
token@cqlsh:default_keyspace> INSERT INTO account (atype, accno, aod, astatus, balance, custid)
... VALUES ('saving', 1001, '2023-02-03', 'active', 5000, 1);
token@cqlsh:default_keyspace> INSERT INTO account (atype, accno, aod, astatus, balance, custid)
... VALUES ('saving', 1002, '2023-03-03', 'active', 10000, 2);
token@cqlsh:default_keyspace> INSERT INTO account (atype, accno, aod, astatus, balance, custid)
... VALUES ('saving', 1004, '2021-12-01', 'notactive', 100, 4);
token@cqlsh:default_keyspace> INSERT INTO account (atype, accno, aod, astatus, balance, custid)
... VALUES ('current', 1002, '2020-01-01', 'active', 1000, 3);
token@cqlsh:default_keyspace> INSERT INTO account (atype, accno, aod, astatus, balance, custid)
... VALUES ('current', 1003, '2020-01-01', 'active', 1000, 3);
token@cqlsh:default_keyspace> SELECT * FROM account;
```

| atype   | accno | aod        | astatus   | balance | custid |
|---------|-------|------------|-----------|---------|--------|
| saving  | 1001  | 2023-02-03 | active    | 5000    | 1      |
| saving  | 1002  | 2023-03-03 | active    | 10000   | 2      |
| saving  | 1004  | 2021-12-01 | notactive | 100     | 4      |
| current | 1002  | 2020-01-01 | active    | 1000    | 3      |
| current | 1003  | 2020-01-01 | active    | 1000    | 3      |

(5 rows)

**3. Display the details of account 1001.**

```
token@cqlsh:default_keyspace> SELECT * FROM account WHERE atype='saving' AND accno=1001;
```

| atype  | accno | aod        | astatus | balance | custid |
|--------|-------|------------|---------|---------|--------|
| saving | 1001  | 2023-02-03 | active  | 5000    | 1      |

(1 rows)

**4. Display account status along with account number of all the saving accounts.**

```
token@cqlsh:default_keyspace> SELECT accno, astatus FROM account WHERE atype='saving';
```

| accno | astatus   |
|-------|-----------|
| 1001  | active    |
| 1002  | active    |
| 1004  | notactive |

(3 rows)

##### 5. Display the highest balance of the savings account.

```
token@cqlsh:default_keyspace> SELECT MAX(balance) FROM account WHERE atype='saving';
```

| system.max(balance) |
|---------------------|
| 10000               |

(1 rows)

##### 6. Create table weblogs as per the following given data types.

(Check the function now(),dateof() for current time and date of the system)

```
token@cqlsh:default_keyspace> create table weblogs (
    ...     page_id uuid,
    ...     page_name text,
    ...     insertion_time timestamp,
    ...     page_count counter,
    ...     PRIMARY KEY((page_id, page_name), insertion_time)
    ... );
token@cqlsh:default_keyspace> SELECT * FROM table;
```

##### 7. Insert records in order to understand the use of all different types.

```
token@cqlsh:default_keyspace> UPDATE weblogs SET page_count = page_count + 1
    ... WHERE page_id = uuid() AND page_name = 'my_page' AND insertion_time = toTimestamp(now());
token@cqlsh:default_keyspace> SELECT * FROM weblogs;
```

| page_id                              | page_name | insertion_time                  | page_count |
|--------------------------------------|-----------|---------------------------------|------------|
| b212f54c-6162-448b-8118-c1dd7b00ed81 | my_page   | 2025-03-30 16:12:09.463000+0000 | 1          |

(1 rows)

8. Let suppose that a user visited the same page for the second time then the page\_count should be incremented to one.write a cassandra query to insert the value in the given table.your query should generate the output as shown in below Figure.

```
token@cqlsh:default_keyspace> UPDATE weblogs SET page_count = page_count + 1
... WHERE page_id = 3dde0197-f452-4dbb-a8d7-67c8c8ec4733 AND page_name = 'my_page'
... AND insertion_time = '2023-09-05 11:09:19.732000+0000';
token@cqlsh:default_keyspace> SELECT * FROM weblogs;
```

| page_id                              | page_name | insertion_time                  | page_count |
|--------------------------------------|-----------|---------------------------------|------------|
| b212f54c-6162-448b-8118-c1dd7b00ed81 | my_page   | 2025-03-30 16:12:09.463000+0000 | 1          |
| 3dde0197-f452-4dbb-a8d7-67c8c8ec4733 | my_page   | 2023-09-05 11:09:19.732000+0000 | 1          |

(2 rows)

**9. Check what if we execute the following query. Justify your answer.**

**select \* from weblogs where page\_name='mypage2'**

```
token@cqlsh:default_keyspace> SELECT * FROM weblogs WHERE page_name='mypage2';
InvalidRequest: Error from server: code=2200 [Invalid query] message="Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING"
```

Justification:

This query will fail because page\_name is part of the composite primary key but is not the partition key. Cassandra requires at least the partition key (page\_id) in the WHERE clause.

**10.Delete one of the Pages from the table.**

```
token@cqlsh:default_keyspace> DELETE FROM weblogs WHERE page_id = 3dde0197-f452-4dbb-a8d7-67c8c8ec4733 AND page_name = 'my_page';
token@cqlsh:default_keyspace> |
```

**11. Insert a few page records and display them like the following figure.  
understand the purpose of partition key column, clustering key column.**

```

token@cqlsh:default_keyspace> DELETE FROM weblogs WHERE page_id = 3dde0197-f452-4dbb-a8d7-67c8c8ec4733 AND page_name = 'my_page';
token@cqlsh:default_keyspace> UPDATE weblogs SET page_count = page_count + 1
... WHERE page_id = uuid() AND page_name = 'my_page3' AND insertion_time = toTimestamp(now());
token@cqlsh:default_keyspace> UPDATE weblogs SET page_count = page_count + 2
... WHERE page_id = uuid() AND page_name = 'my_page' AND insertion_time = toTimestamp(now());
token@cqlsh:default_keyspace>
token@cqlsh:default_keyspace> UPDATE weblogs SET page_count = page_count + 1
... WHERE page_id = uuid() AND page_name = 'my_page2' AND insertion_time = toTimestamp(now());
token@cqlsh:default_keyspace> UPDATE weblogs SET page_count = page_count + 1
... WHERE page_id = uuid() AND page_name = 'my_page4' AND insertion_time = toTimestamp(now());
token@cqlsh:default_keyspace> UPDATE weblogs SET page_count = page_count + 1
... WHERE page_id = uuid() AND page_name = 'my_page6' AND insertion_time = toTimestamp(now());
token@cqlsh:default_keyspace>
token@cqlsh:default_keyspace> UPDATE weblogs SET page_count = page_count + 1
... WHERE page_id = uuid() AND page_name = 'my_page5' AND insertion_time = toTimestamp(now());
token@cqlsh:default_keyspace> SELECT * FROM weblogs;

```

| page_id                              | page_name | insertion_time                  | page_count |
|--------------------------------------|-----------|---------------------------------|------------|
| fc79a0bc-99ff-4e83-b50f-7b041cb06705 | my_page6  | 2025-03-30 16:18:49.333000+0000 | 1          |
| b0516af9-4632-4f51-bcbf-adea0dcee941 | my_page5  | 2025-03-30 16:19:02.513000+0000 | 1          |
| 183bb788-9179-460d-93f2-90e179e34194 | my_page2  | 2025-03-30 16:18:16.869000+0000 | 1          |
| b212f54c-6162-448b-8118-c1dd7b00ed81 | my_page   | 2025-03-30 16:12:09.463000+0000 | 1          |
| 1242abba-d33a-4558-a93b-b5bc0920323c | my_page   | 2025-03-30 16:18:03.333000+0000 | 2          |
| 3caad24e-6929-4a2b-b653-0b25b04cda0f | my_page3  | 2025-03-30 16:17:47.503000+0000 | 1          |
| 396fb99b-a7ad-4a9c-b462-0ea59f6e6971 | my_page4  | 2025-03-30 16:18:35.437000+0000 | 1          |

(7 rows)

## 12.Display the total of all page record count.

```

token@cqlsh:default_keyspace> SELECT SUM(page_count) FROM weblogs;

```

```

system.sum(page_count)
-----

```

8

(1 rows)

Warnings :

Aggregation query used without partition key

## 13.Display how many pages are stored in the table.

```

token@cqlsh:default_keyspace> SELECT COUNT(*) FROM weblogs;

```

```

count
-----

```

7

(1 rows)

Warnings :

Aggregation query used without partition key

## 14.Display the page which has the highest number of visitors(page\_count).

```
token@cqlsh:default_keyspace> SELECT page_name, MAX(page_count) FROM weblogs;

page_name | system.max(page_count)
-----+-----
my_page6 | 2

(1 rows)

Warnings :
Aggregation query used without partition key
```

### 15. Explore, Are Joins and Subqueries supported in Cassandra? If Yes, then How? If No, then Why?

No, Joins and Subqueries are not supported in Cassandra.

#### Why?

- **Denormalization is encouraged:** Cassandra is a NoSQL database optimized for speed and scalability, so it avoids joins to improve performance.
- **No Complex Queries:** Cassandra does not have relational integrity constraints like foreign keys.
- **Designed for Distributed Systems:** Since data is stored across multiple nodes, performing joins would be inefficient and slow.
- **Alternative Approach:** Data should be pre-joined and stored in separate tables based on access patterns.

#### Summarised learning:

In this lab, I explored key concepts and operations in Cassandra, including table creation, data insertion using upsert operations, and querying. I worked with partition and clustering keys to organize data, and explored counters for managing increment/decrement operations on values. Additionally, I gained hands-on experience with various query operations such as retrieving specific data, calculating aggregations, and managing counter columns. I also worked with more advanced topics like timestamp and counter-based tables, and understood the importance of partition and clustering keys in efficient data management.